


Why a GARP-blind null outperforms the real repair, on the sample's largest-dose trace (qwen2.5:1.5b, reciprocal, replicate 12; $T = 21$ valid rounds; one illustrative trace, not a statistical result)



C. Where each operator spends the identical L_1 budget of 111.64

the repair puts all of it on 5 of the 21 rounds; the null spreads it over all 21 ($\lambda = 0.092$)

