
What Does Repairing Choice Inconsistency Actually Buy?

A Budget-Set Diagnosis

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Repairing an AI agent’s incoherent preferences is not new: inference-time transi-
2 tivity repairs, isotonic projections onto preference-defined cones, and training-time
3 rationality penalties have all been proposed, several with reported downstream
4 gains. None controls for displacement magnitude: no result establishes whether
5 restoring coherence itself, not mere displacement toward an interior point, is what
6 any reported gain comes from. We build that control. Given an LLM agent’s own
7 choices over budget sets, we compute the minimal MILP perturbation restoring
8 GARP-consistency, then compare it against a GARP-blind null operator spending
9 the identical displacement budget shrinking every bundle toward a fixed center.
10 Under the paper’s original exogenous payoff, the null outperforms the real repair by
11 more than 2x (mean payoff gain 0.0220 vs. 0.0091, Wilcoxon $p = 3.98 \times 10^{-10}$,
12 winning 81% of traces). A second, independently designed payoff built to re-
13 move the first one’s exploitable fixed target reproduces the same verdict (null
14 0.0217 vs. real 0.0072, $p = 1.24 \times 10^{-5}$, 70.8% win rate); we decline a third as
15 payoff-shopping. Trace extremity predicts the null’s advantage in both experiments.
16 The paper’s standing contributions are instead a capacity-deconfounded identifica-
17 tion strategy for the coherence–competence question, which remains open, and a
18 discard-selection instrument-validity finding.

19 1 Introduction

20 An LLM agent asked to allocate a fixed budget across two goods at varying prices, repeatedly, will
21 frequently violate the Generalized Axiom of Revealed Preference (GARP; Afriat, 1967, Varian, 1982):
22 it prefers bundle x_1 to x_2 at one budget line and the reverse at another, in a way no monotone utility
23 function can rationalize. The same budget-allocation instrument is the standard paradigm for testing
24 GARP experimentally on human subjects [Andreoni and Miller, 2002, Choi et al., 2007, 2014], and a
25 growing literature applies it to GPT-family models directly, without a repair step, to ask whether they
26 satisfy classical rationality axioms at all [Chen et al., 2023]. A separate, growing literature repairs
27 exactly this kind of inconsistency in LLM systems — restoring transitivity in pairwise judgments,
28 projecting cardinal scores onto monotone cones, penalizing incoherence during training — and
29 several of these systems report that the repaired outputs are downstream-better than the raw ones.
30 What none of them establishes is a *dose*: how much better, as a function of how much repair was
31 needed, isolated from the capacity and training confounds that come from comparing a richer model
32 class to a poorer one.

33 We study this question in a setting where it can be answered cleanly. An LLM agent makes a
34 sequence of budget-constrained choices over two goods. We compute, as a single mixed-integer
35 linear program, the minimal quantity perturbation that would make the agent’s own realized choice

sequence GARP-consistent — a projection applied *post hoc* to a frozen agent, so nothing about its weights, training, or decoding changes across doses. The size of that perturbation is a graded, cardinal dose. We score both the raw and the repaired sequence against a fixed, equal-weight Cobb–Douglas payoff whose optimum at each budget line is a closed-form function of prices and income alone — never of the agent’s own choices — so the outcome measure is exogenous to the preference data by construction, not a restatement of coherence. Tracing dose against Δ payoff gives the relationship this paper reports.

That relationship turns out not to isolate what it was built to isolate. A control absent from every published axiom-enforcement result we are aware of — an operator that knows nothing about GARP but spends the identical displacement budget shrinking every bundle toward a fixed interior point — outperforms the real GARP-restoring repair on the exogenous payoff, significantly, at both model scales, under two independently designed payoffs (§5.1). This is the paper’s central finding: restoring rationality, as originally proposed, does not buy more exogenous payoff than mere displacement toward the interior of the budget line does on its own. We report this as a negative result reached twice by construction, not as a null we happened to observe once.

We do not claim any of the individual pieces are new; §2 concedes exactly what is not. What has not been done is the conjunction: an agent’s own realized choices, a graded coherence-indexed dose, and an outcome that is not itself a preference judgment, run together so that the identification is clean. We run this at a model scale where a pilot study found measurable coherence headroom (1.5B parameters), and repeat the identical pipeline at a scale with almost none (3B parameters) as an identification control on the method itself, not as a second point on a shared dose curve across scale.

Two further results were not predicted by the design. The reciprocal-price framing manipulation intended to induce coherence variation at 1.5B produced, in a pilot run, a large apparent effect that disappears once discard-selection is corrected with a capped retry protocol: the pilot’s naive estimate and the corrected main-experiment estimate are both statistically indistinguishable from zero, and the two should not be read as a sign reversal of a real effect. We report the discard rate itself, not the CCEI point estimate, as the more reliable signal that the manipulation disrupted the agent (our claim C3). Wang et al. [2025a] already recommend reporting invalid-response proportions as a checklist item; our contribution is a quantitative demonstration of how much a naive discard rule can distort a measured effect at this specific scale, broken down by attempt outcome in §5.2, not the recommendation itself. Separately, splitting single-turn elicitation into separate sequential calls at 1.5B produces a large GARP-pass-rate collapse, moving in the *opposite* direction from an independently published single-turn-vs-multi-turn finding in the same domain and model family at a larger scale, with no discard confound on either arm.

Contributions. (1) An application — not a novel method — of Demuyneck and Rehbeck [2023]’s minimal-quantity-error MILP projection to an agent’s own realized choices, independently verified on every output rather than trusted from solver status alone. (2) A closed-form exogenous payoff, fixed before any data collection, audited adversarially for leakage and for the mechanical confound that larger violations simply have more room to improve. (3) A distance-matched, GARP-blind null-operator control — absent from every published axiom-enforcement result we are aware of — run against two independently designed exogenous payoffs, showing that the measured raw dose–response relationship is not evidence that GARP-restoration specifically, rather than mere displacement toward the interior, is what buys the exogenous payoff gain (§5.1). (4) An instrument-validity finding (discard-selection bias under a disruptive manipulation), corrected for by a stated retry protocol and reported per cell rather than pooled away. (5) A single-turn-vs-multi-turn elicitation-format manipulation that produces a large GARP-pass-rate collapse at 1.5B, in the *opposite* direction from an independently published single-turn-vs-multi-turn finding in the same budget-allocation domain and the same model family (Qwen2.5) at a larger scale.

2 Related Work

Repairing an AI system’s incoherent preferences is not new; Table 1 positions nine published systems against the four criteria that jointly define our claim: an agent’s own choices (not a judgment about third-party items), an exogenous payoff (no preference judgment enters it), a graded rather than binary dose, and a formal minimum-distance projection. None occupies all four.

Table 1: Where prior repair systems sit. *Own choices*: the repaired object is a choice the agent made for itself from a budget set, not a judgment about third-party items. *Exogenous payoff*: an outcome into which no preference judgment enters.

	Own choices	Exogenous payoff	Graded dose	Min.-distance projection
POISE [Wang et al., 2026]	no	no	no	yes
Rationality layer [Chadwick et al., 2025]	no	no	no	yes
TrustJudge [Wang et al., 2025b]	no	no	no	no
CONSISTRE [Sun, 2026]	no	partial	no	no
LLM-RankFusion [Zeng et al., 2024]	no	yes	no	no (declined)
Innate preferences [Buchanan and Foster, 2026]	yes	no	no	no
GARP-EFM [Aguar and Kashaev, 2026]	no	yes	no	no
HRC/DSPPO [Huang et al., 2026]	no	no	yes	no
Control vectors [Cook et al., 2026]	yes	no	yes	no
This paper	yes	yes	yes	yes

89 **POISE** [Wang et al., 2026] is the sharpest vocabulary collision: pool-adjacent-violators projection
90 onto a closed convex chain-monotone cone, with a proved Pythagorean guarantee the edit lands weakly
91 closer to a posited ground truth. We concede the minimum-distance priority unqualified: projection
92 onto a closed convex set is non-expansive in L_2 , licensing that guarantee; the GARP-consistent set is
93 a union of polyhedra, hence not convex, so no analogue transfers here (§3.2). POISE also projects
94 a *teacher’s* offline labels, not an agent’s own choices. Three further inference-time repairs, each
95 on a different object: Chadwick et al. [2025] (synthetic election data, no downstream comparison),
96 TrustJudge [Wang et al., 2025b] (third-party judgments, no degree parameter), CONSISTRE [Sun,
97 2026] (relation-extraction triples, no Afriat machinery). LLM-RankFusion [Zeng et al., 2024] scores
98 its rankings against real TREC relevance labels — the only one of these four with a genuinely
99 exogenous payoff — but explicitly declines a minimum-distance rank-aggregation method, naming
100 Kemeny aggregation and rejecting it as NP-hard in favor of simple Borda-count voting across rankers.
101 One training-time system alters the agent itself: Buchanan and Foster [2026] (IIA-invariance loss,
102 no downstream evaluation). A further training-time system is not an LLM agent at all: Aguar and
103 Kashaev [2026] fine-tune Chronos-2, a time-series forecaster, on synthetic utility-maximizing-agent
104 data to predict human consumers’ own choices, with a downstream evaluation (17–31% bundle-
105 prediction error reduction) neither of the other two above reports.

106 A parallel line varies the transitivity of the *preference model* rather than an agent’s choices: GPM
107 [Zhang et al., 2025] and HRC/DSPPO [Huang et al., 2026] both make the cycle-tolerant arm a strict
108 superset model class, confounding coherence with capacity by construction, and on GPM’s own
109 length-controlled metric that arm *loses* in 18 of 24 cells. HRC/DSPPO already publishes its own
110 inverted-U dose–response curve over a training-time schedule weight — we say so before claiming
111 anything about grading — but the dose weights a learned third-party proxy during training, not an
112 agent’s own realized choices, and the outcome is an LLM-judge score throughout, where ours acts
113 post hoc and is exogenous. Our closest theoretical neighbour, Andrews [2026], proposes 1 – CCEI
114 as a training penalty and states plainly, repeatedly, that coherence is *not* sufficient for good behaviour
115 — argued a priori, never measured, never asking whether *imposing* coherence helps or hurts; ours is
116 the empirical counterpart to a proposal that has circulated unrun.

117 **The sharpest objection is psychometric, and it is in PNAS.** Nitsch et al. [2022] find CCEI/Houtman–
118 Maks [Houtman and Maks, 1985] reliability never reaches acceptable levels across eight datasets,
119 and that participants given the chance to revise their own inconsistent choices saw mean CCEI *fall*.
120 That revision arm is not a repair operator (a random subset of choices, no consistency objective,
121 no guarantee of fewer violations) — ours attains rationalizability by construction and verifies it
122 independently, the same distinction that disposes of Yamin et al. [2026a]’s isotonic-calibration repair
123 (worse in 14 of 16 cells, projecting onto a different set than the one it is graded on). A companion
124 paper by an overlapping author team fits revealed-preference-derived cost functions to LLM decisions
125 to diagnose alignment and steerability rather than to repair anything [Yamin et al., 2026b] — the same
126 GARP/Afriat machinery used here for measurement rather than projection. The reliability finding
127 is diagnosed by its own authors as a between-subject-variance problem whose prescribed remedy
128 — a between-groups design — is this paper’s design; the concession we do not shed is that three

independently published axiom-enforcement results (adding Zhu et al. [2025]’s) all point the same adverse way, and a null result here would be a fourth confirmation rather than a discovery, which is why we carry a distance-matched null-operator control neither published negative had.

The closest neighbour on the economics side is invisible to any arXiv sweep. Cook et al. [2026] steer economic choices toward payoff-maximising behaviour via personas, prompt masking, and control vectors — occupying two of our three legs without Afriat machinery or a minimum-perturbation objective; their finding that reframing moves models while persona prompting does not corroborates our manipulation choice, anchored in Wang et al. [2025a]. Wang et al. find *format*, not persona or temperature, moves the Afriat index in the same budget-allocation domain we study, on Qwen2.5-7B-Instruct among other models — the same model family as our 1.5B, at a larger scale: collapsing their multi-turn baseline to single-turn drops CCEI by up to 0.241. Our arm runs the opposite manipulation (single-turn baseline split into multi-turn) and finds the opposite direction: GARP pass rate collapses from 0.40 to 0.10 ($p = 0.0073$; survives Benjamini–Hochberg but not Holm correction across our full test family) with zero discards. This is not a replication: in the same domain and model family, which format is more coherence-degrading reverses between their 7B model and our 1.5B model. We report the disagreement rather than the confirmation we expected.

3 Method

3.1 The projection: Demuynck & Rehbeck’s (2023) minimal-quantity-error MILP, applied and independently verified

This section applies, rather than proposes, a projection method. The MILP formulation below is Demuynck and Rehbeck [2023]’s multiplier-free ordinal characterization of minimal-quantity-error GARP repair, adopted here on an LLM agent’s own choice sequence and verified independently on every output; nothing about the formulation itself is new to this paper.

For a sequence of T observations $(p_t, x_t)_{t=1}^T$, K goods, GARP holds if there is no sequence of choices such that x_{t_1} is revealed preferred to x_{t_2} , ..., x_{t_n} is revealed preferred to x_{t_1} , with at least one strict revealed preference in the cycle. Given a GARP-violating sequence, we seek the minimal perturbation $\tilde{x}_t$ that restores GARP-consistency. The naive formulation parameterizes this with Afriat multipliers, which makes the resulting constraints bilinear in the unknowns once $\tilde{x}$ is itself a decision variable — fixing a candidate preference ordering does not remove the bilinearity, since a price-weighted multiplier term remains a product of two unknowns regardless of ordering. We instead use the multiplier-free ordinal characterization of Demuynck and Rehbeck [2023], under which every constraint is linear jointly in the perturbed bundles, a set of ordinal utility levels $u_t \in [0, 1]$ that carry no multipliers, and binary comparison indicators $U_{t,v} \in \{0, 1\}$ for $t \neq v$. At $T = 25$ this is 600 binaries and 125 continuous variables per trace, solved as a single MILP on the HiGHS backend [Huangfu and Hall, 2018], with no outer search over preference orderings and no alternating scheme.

Three method commitments, each carrying an explicit risk. **Budget exhaustion is imposed as an equality**, $p_t \cdot \tilde{x}_t = I_t$ for every observation (income fixed at 100 throughout this study); this is standard for the design, and it makes the big- M constant in the ordering constraints computable a priori as $\alpha > \max_t I_t$. **A fixed strict-preference margin** $\gamma > 0$, set to $10^{-4} \cdot \min_t I_t$, converts an unattained infimum into an attained minimum: the GARP-consistent set is not closed, so an unregularized projection can read a distance of exactly zero on genuinely violating data. We verified the reported distance is stable across $\gamma \in \{10^{-2}, \dots, 10^{-6}\}$ before treating it as reportable. **Every returned $\tilde{x}$ is verified independently** via the same combinatorial Warshall-closure [Warshall, 1962] GARP check used everywhere else in this paper, never trusted from solver optimality status alone — all 85 GARP-violating traces’ projections in the main experiment were independently re-verified GARP-consistent, with MIP gap $\leq 8.1 \times 10^{-5}$ on every solve.

The objective is L_1 : $\min \sum_t \sum_k w_{t,k} d_{t,k}$ subject to $\tilde{x}_{t,k} - x_{t,k} \leq d_{t,k}$ and $x_{t,k} - \tilde{x}_{t,k} \leq d_{t,k}$, with uniform weights $w = 1$. This keeps the program a pure MILP; L_∞ distance is recorded alongside L_1 for every trace but is not separately analyzed in this paper. L_2 , the literature’s own least-squares index, would require an MIQP solver we do not have configured and is not computed. The *dose* for a trace is its L_1 projection distance $\sum_t \sum_k |x_{t,k} - \tilde{x}_{t,k}|$, comparable across sessions without further normalization because income is fixed throughout — the same style of graded, minimum-cost severity measure as Dean and Martin [2016], applied here to an LLM’s own choices rather than lab or survey human data and paired with an exogenous behavioral payoff rather than left as a standalone

183 index. A Cobb–Douglas demand share-fitted to each trace’s own observed data is computed as a
 184 feasibility incumbent and sanity ceiling on the reported distance, recorded only as an upper-bound
 185 diagnostic; Appendix A(1) checks in detail that it never reaches the solver and so never influences
 186 any reported projection.

187 3.2 What can be guaranteed, and what cannot: positioning against convex projection

188 §2 positions this projection against the closest published structural analogue, a proved-optimal
 189 projection onto a convex monotone cone [Wang et al., 2026]: because the GARP-consistent set is
 190 a union of polyhedra rather than a single convex set, the non-expansiveness that licenses a weakly-
 191 closer-to-ground-truth guarantee there has no analogue here. Concretely for this formulation, the
 192 prior operator’s cone is fixed by an ordering supplied as input, whereas a GARP repair must decide
 193 which revealed-preference comparisons to give up — the ordering search is absorbed into the
 194 binary comparison indicators of §3.1 rather than eliminated, which is why no distance-minimization
 195 guarantee analogous to the convex case is claimed.

196 3.3 The exogenous payoff

197 The payoff must not be derived from the agent’s own revealed choices — a utility function fit to
 198 the agent’s choices (such as the projection’s own feasibility incumbent above) would be circular if
 199 used as the outcome measure, rewarding the agent for resembling its own revealed preferences rather
 200 than measuring anything external to them. We instead fix, before any model was queried and never
 201 re-estimated from any agent’s choices, an equal-weight Cobb–Douglas valuation $U_{\text{exo}}(x) = x_A^{0.5} x_B^{0.5}$.
 202 At budget line (p_t, I_t) the optimal bundle x_t^* is the closed-form Cobb–Douglas demand $x_{t,A}^* =$
 203 $0.5I_t/p_{t,A}$, $x_{t,B}^* = 0.5I_t/p_{t,B}$ — a function of prices and income alone, never of the agent’s observed
 204 or projected bundle. The payoff score is the efficiency ratio $\text{payoff}(x_t) = U_{\text{exo}}(x_t)/U_{\text{exo}}(x_t^*) \in$
 205 $(0, 1]$, computed identically for the raw and projected sequence at the same budget lines, so that
 206 $\Delta\text{payoff} = \text{payoff}(\tilde{x}) - \text{payoff}(x)$ is comparable across conditions, models, and repairs. This
 207 mirrors — not the specifics, but the structure of — a money-metric utility index, the standard device
 208 for welfare comparison over the same GARP/Afriat objects this paper’s method already relies on. A
 209 genuinely exogenous, real-world payoff of this kind is rare in this literature: Ouyang et al. [2024]
 210 score LLM-generated capital-expenditure forecasts against firms’ realized future spending and find the
 211 relationship between alignment intensity and forecast accuracy is *non-monotonic* — light alignment
 212 helps, over-alignment hurts — but without a placebo intervention of matched size, that design cannot
 213 separate an alignment-specific effect from a generic effect of any fine-tuning intervention, the same
 214 separation our dose–response design is built to make.

215 We audited this design adversarially, before treating any dose–response result as reportable, for
 216 four failure modes: a shared-data leak between payoff and projection, a mechanical link between
 217 GARP-consistency and payoff, a projection direction secretly aimed at the payoff optimum, and the
 218 confound that larger-dose traces simply start worse and have more room to improve. None survived
 219 the audit (Appendix A gives every check and number); the raw dose–response statistic we report in
 220 §5.1 is not an artifact of any of the four. A fifth failure mode — that this payoff’s own concavity
 221 could reward mere displacement without any GARP-specific effect — required a new control rather
 222 than an audit of the existing design, and is addressed separately in §3.4.

223 3.4 Two robustness controls against a payoff-geometry confound

224 The payoff of §3.3 reduces algebraically to $\text{payoff}(s) = 2\sqrt{s(1-s)}$, where s is the expenditure
 225 share on good A: concave, with a single interior maximum at $s = 0.5$ identical across every trace.
 226 Any operator that moves a bundle toward that one fixed point raises its payoff, whether or not the
 227 movement has anything to do with restoring GARP-consistency. To test whether the dose–response
 228 relationship of §5.1 reflects GARP-restoration specifically rather than this geometric fact, we construct
 229 a **null operator**: for each GARP-violating trace, shrink every observed bundle toward the exogenous
 230 optimum x_t^* by a single shared factor $\lambda \in [0, 1]$, chosen so the null’s total L_1 displacement exactly
 231 matches the real projection’s own dose on that trace. The null never calls the GARP check and does
 232 not, in general, restore consistency; it is scored by the same payoff function used everywhere else.

233 Because the first payoff’s exploitable property is a single fixed target identical across every trace, we
 234 build a second, independently designed payoff that removes it: a per-trace Cobb–Douglas weight

235 $\alpha_s \sim \text{Uniform}(0.05, 0.95)$, drawn independently for each of the 85 traces and for each of $K = 20$
 236 replicate draws, read only at scoring time and never by the projection or the null construction. Under
 237 this second payoff we report two null constructions, never pooled: a **primary**, information-fair null
 238 (the identical construction above, unaware of α_s) and an **oracle** null (aims directly at the true per-trace
 239 optimum $x_{\alpha_s}^*$), reported only as an upper bound since it uses information no real repair algorithm has.

240 4 Experimental design

241 4.1 Models and conditions

242 Two locally-hosted open-weight models, run entirely on local compute at zero API cost:
 243 qwen2.5:1.5b-instruct (the headroom model, per a pilot study that found measurable within-
 244 condition coherence variation at this scale) and llama3.2:3b-instruct (the null-effect control, per
 245 the same pilot finding almost no coherence headroom at this scale — mean baseline CCEI 0.99). A
 246 third model family was tested for feasibility but excluded after a reproducible multi-model-residency
 247 crash under sustained rotation on the local machine; the two-model, model-major design (all sessions
 248 for one model complete before the next model loads) is both a wall-clock optimization and, after
 249 that observed failure, a correctness requirement. **Compute.** All model inference ran locally via
 250 4-bit-quantized weights on a single consumer CPU-class machine, no GPU and no commercial API
 251 calls; the full $N = 30$ design across both models and every condition completed in approximately
 252 2.1 hours wall-clock. Every MILP projection (§3.1) solved in under 5 seconds on the same machine,
 253 CPU-only, with MIP gap $\leq 8.1 \times 10^{-5}$ on every solve.

254 Three conditions at 1.5B, two at 3B. **Baseline:** direct per-token-return framing, single turn, all 25
 255 rounds in one prompt and one response. **Reciprocal:** identical single-turn format with an inverted-
 256 price framing manipulation, the strongest manipulation identified in the pilot. **Multiturn** (1.5B only):
 257 baseline framing, but each of the 25 rounds is a separate sequential API call carrying the running
 258 conversation history, isolating the elicitation-format mechanism from the price-framing mechanism;
 259 the arm is motivated by an independently published finding, in the same budget-allocation domain
 260 we study, that collapsing multi-turn elicitation to single-turn moves CCEI by up to -0.241 for a
 261 same-family model at a larger scale [Wang et al., 2025a], an order of magnitude larger than this
 262 paper’s own pilot framing effect. Our arm tests the manipulation in the opposite direction (single-turn
 263 baseline split into multi-turn), since that is the direction our own baseline format already uses. It is
 264 not added at 3B, whose role is a null-effect control on the identification strategy rather than a second
 265 scale point requiring the full condition set.

266 4.2 Budget sets and power

267 Each of $N = 30$ independent replicates per (model, condition) cell draws its own fresh $T = 25$, $K =$
 268 2 budget set (continuous-density prices, $M, N \sim U[0.1, 1.0]$ i.i.d. rejected unless $\max(M, N) \geq 0.5$,
 269 budget exhaustion enforced), seeded independently from every other replicate — unlike a pilot
 270 study’s single shared set, needed because the main experiment targets independent between-condition
 271 replication, not test–retest reliability. Continuous, independently-drawn prices with enforced budget
 272 exhaustion also avoid a known pathology: CCEI can read exactly 1.0 on data that violates GARP
 273 whenever a comparison lands on exact budget-line equality, an event common under grid-sampled
 274 or rounded prices [Andrews, 2026]. $N = 30$ gives power ≥ 0.999 against the pilot’s full framing-
 275 effect magnitude on the pass-rate metric and ≥ 0.80 down to $\sim 60\%$ attenuation, but is explicitly
 276 underpowered for a CCEI-scale effect of the pilot’s magnitude at 1.5B (minimum detectable ≈ 0.11),
 277 which we report rather than paper over. Every replicate’s Bronars power [Bronars, 1987] is recorded;
 278 all 150 draws held power ≥ 0.998 .

279 4.3 The discard-selection problem, and its correction as a stated contribution

280 A pilot run found that under reciprocal framing at 1.5B, 52% of sessions (13 of 25) failed the required
 281 output-format contract and were silently discarded under the pilot’s own naive handling — standard
 282 practice, and one this paper measures the consequences of directly rather than assumes. We treat this
 283 as a claim about instrument validity in its own right, not a footnote to the dose–response result, true
 284 regardless of whether projection helps, hurts, or does nothing downstream. The main experiment
 285 implements a capped retry protocol: a session failing the $\geq 20/25$ -valid-rounds threshold is retried

Table 2: Headline results per (model, condition) cell. GARP pass and CCEI are computed on kept traces only; dose and Δ payoff are computed on the subset of kept traces that violate GARP.

Model	Condition	n kept	mean CCEI	GARP pass	mean dose (L_1)	mean Δ payoff
llama3.2:3b	baseline	30	0.9900	0.73	5.01	+0.0018
llama3.2:3b	reciprocal	28	0.9713	0.39	6.27	+0.0016
qwen2.5:1.5b	baseline	30	0.9522	0.40	16.68	+0.0090
qwen2.5:1.5b	multiturn	30	0.9540	0.10	14.54	+0.0083
qwen2.5:1.5b	reciprocal	24	0.9413	0.38	12.02	+0.0065

up to three total attempts with a fresh seed offset, every attempt recorded; a slot still failing after three attempts is a residual discard, excluded from CCEI/GARP but retained for audit.

Broken down by attempt outcome at 1.5B reciprocal framing, retry-rescued sessions ($n = 7$, recovered on attempt 2 or 3) are not a close stand-in for first-attempt sessions ($n = 17$): they show *higher* mean CCEI (0.965 vs. 0.932) but a substantially *lower* GARP pass rate (14% vs. 47%). The sessions that never produce a usable trace even after three attempts ($n = 6$, three with zero valid rounds recorded at all) show the lowest CCEI among the three groups on the handful that can be scored (0.882, $n = 3$) — consistent with, though not conclusive given the small residual sample, some of the same selection concern the pilot’s naive handling raised persisting at a smaller scale even after correction; the retry protocol reduces but does not eliminate it. We report first-attempt and residual post-retry discard rates as a first-class per-condition outcome for every cell.

5 Results

187 attempt-records were collected across 150 replicate slots (2 models $\times$ up to 3 conditions $\times$ 30 replicates); 142 slots kept a usable trace, 8 were residual discards after three attempts. Every cell reached its full 30 replicates.

5.1 The dose–response relationship, and the control that overturns it (C1)

Across all 85 GARP-violating traces with a computed, independently-verified projection (60 at 1.5B, 25 at 3B): Spearman $\rho = 0.729$ ($p < 0.0001$), Pearson $r = 0.821$ ($p < 0.0001$), mean Δ payoff +0.0091 (one-sample $t = 5.41$, $p < 0.00001$), with 70 of 85 traces (82%) improving after repair and 15 (18%) worsening. Figure 1 shows the relationship split by model on shared axes. **Both scales show a significant, positive relationship, and it is stronger at the headroom model:** $\rho = 0.756$ ($p < 0.0001$, $n = 60$) at 1.5B versus $\rho = 0.614$ ($p = 0.0011$, $n = 25$) at 3B. We tested the mechanical confound that larger-dose traces might simply start from worse raw payoff and therefore have more room to improve. The confound is real (dose vs. raw payoff, Pearson $r = -0.41$, $p < 0.0001$) but does not explain the relationship away: the partial correlation of dose against Δ payoff controlling for raw payoff is $r = 0.784$ ($p = 7 \times 10^{-19}$), and an OLS regression of Δ payoff on dose and raw payoff jointly finds dose highly significant ($t = 11.4$, $p < 10^{-16}$) with raw payoff’s own marginal contribution not distinguishable from zero ($p = 0.18$) once dose is included. Full detail in Appendix A.

This is where a paper that stopped at the raw correlation would end, and it is not where this one ends. The confound just tested is a confound on the *severity* of the raw violation, not on the specific *geometry* of the payoff function — it says nothing about whether restoring GARP-consistency *specifically*, rather than merely displacing choices toward some interior point, is what buys the payoff gain above. §3.4 builds the control that closes this gap: a GARP-blind null operator, spending the identical L_1 displacement budget as the real repair, that does nothing but shrink every bundle toward the exogenous payoff’s own fixed optimum.

Experiment 1 (original payoff). The null operator outperforms the real repair by roughly $2.4\times$: mean Δ payoff_{null} = 0.0220 versus mean Δ payoff_{real} = 0.0091, Wilcoxon signed-rank $p = 3.98 \times 10^{-10}$, with the null winning 69 of 85 traces (81%). This holds independently at both models. The partial correlation of dose against Δ payoff_{real}, controlling for the null’s own mechanically-predicted gain, remains positive and significant ($r = 0.574$, $p = 1.13 \times 10^{-8}$) — dose retains real information beyond what the null predicts — but the paired comparison is unambiguous: at the identical displacement

Projection dose (L_1) vs. change in exogenous payoff, by model
(each point: one GARP-violating trace, post hoc projection; shared axes both panels)




Figure 1: Projection dose (L_1) against the change in exogenous payoff after repair, for every GARP-violating trace, split by model on shared axes so the two panels’ slopes are directly comparable. Both relationships are positive and significant; the headroom model’s (left) is steeper and tighter than the null-control’s (right), though the control’s relationship is itself significant rather than null.

328 budget, an operator with no knowledge of GARP buys more than twice the payoff the real repair
329 does.

330 **Experiment 2 (corrected payoff).** The original payoff’s exploit is a single fixed target ($s = 0.5$)
331 identical across every trace; §3.4 removes this by drawing an independent Cobb–Douglas weight
332 $\alpha_s \sim \text{Uniform}(0.05, 0.95)$ per trace, repeated across $K = 20$ independent draws for robustness. The
333 qualitative result is unchanged: the primary, information-fair null still outperforms the real repair
334 (mean $\Delta\text{payoff}_{\text{null}} = 0.02167$ versus mean $\Delta\text{payoff}_{\text{real}} = 0.00723$, Wilcoxon $p = 1.24 \times 10^{-5}$, win
335 rate 70.8%), significant in 20 of 20 draws. An oracle null privileged to know the true per-trace target
336 does better still (mean $\Delta\text{payoff}_{\text{null}} = 0.02224$, $p = 1.35 \times 10^{-9}$, win rate 80.4%, significant in 20/20
337 draws), reported for completeness rather than as the primary comparison. The partial correlation
338 of dose with $\Delta\text{payoff}_{\text{real}}$, controlling for the null, attenuates from Experiment 1’s $r = 0.574$ to a
339 mean $r \approx 0.37$ across the 20 draws, significant in 14–15 of 20 — a real but more modest signal, not
340 distinguishable from the draw-to-draw noise the 20 draws themselves already show.

341 **No third payoff was attempted, and this is a stated scope decision, not a resource constraint.**
342 Two independently designed, independently motivated payoffs — one the paper’s original, one built
343 specifically to remove the first one’s known exploitable structure — converging on the same negative
344 paired-comparison result is the finding. Redesigning the payoff a third time only because the second
345 attempt also came back negative would be payoff-shopping: iterating the yardstick until GARP-
346 restoration happens to look good, exactly the researcher-degree-of-freedom problem an exogenous
347 payoff was introduced to rule out.

348 **A partial mechanistic note.** Trace extremity (1 — raw payoff under the original fixed payoff, a
349 stable measure of the raw bundles independent of which payoff later scores them) predicts the null’s
350 advantage strongly in both experiments (Experiment 1: Pearson $r = 0.679$; Experiment 2 primary
351 null: mean $r = 0.631$ across 20 draws), more strongly than it predicts the real repair’s own gain
352 in either. This pattern alone does not cleanly separate “centering” from a more general “moving
353 an extreme trace toward the interior helps, regardless of exactly where” account, since Experiment
354 2’s primary null is mechanically the same fixed-center operator as Experiment 1’s. Experiment 2’s
355 oracle null, which targets a genuinely different, per-trace-varying point, shows essentially the same
356 extremity-advantage (mean $r = 0.624$) — read as partial, not conclusive, support for an “escape a
357 bad start” mechanism rather than pure centering, offered as a discussion-level hypothesis rather than
358 a load-bearing result with its own significance test. Figure 2 illustrates the same idea concretely on
359 a single trace: the rounds that generate the null’s advantage there are precisely the extreme-share
360 rounds the real repair leaves untouched.

361 **C1 as originally proposed — that restoring GARP-consistency specifically, beyond mere dis-**
362 **placement toward the interior of the budget line, buys more exogenous payoff — is not supported**
363 **under either payoff design.** The raw dose–response relationship reported at the top of this section is
364 real; it is not evidence for the specific mechanism C1 claims.

Why a GARP-blind null outperforms the real repair, on the sample's largest-dose trace
(qwen2.5:1.5b, reciprocal, replicate 12; $T = 21$ valid rounds; one illustrative trace, not a statistical result)




Figure 2: Why the size-matched, GARP-blind null outperforms the real repair, shown on one trace rather than argued from aggregate statistics. With $K = 2$ and budget exhaustion the payoff is exactly a function of the expenditure share s on good A: $2\sqrt{s(1-s)}$ under the original design (panel A, peak fixed at $s = 0.5$ for every trace) and $s^{\alpha_s}(1-s)^{1-\alpha_s}/[\alpha_s^{\alpha_s}(1-\alpha_s)^{1-\alpha_s}]$ under the corrected one (panel B, peak at this trace’s own draw $\alpha_s = 0.704$, draw $k = 0$ of the 20). The trace is qwen2.5:1.5b, reciprocal, replicate 12 — the largest-dose violation in the sample ($L_1 = 111.64$, $T = 21$ valid rounds), hand-audited in Appendix A; every plotted quantity is recomputed by resolving the projection MILP. Panel C gives the mechanism: the minimal-perturbation repair spends its whole displacement budget on 5 of the 21 rounds (42% on round 3 alone) and leaves the other 16 untouched, where the null spreads the identical budget over all 21 ($\lambda = 0.092$). *On the rounds it does touch the repair beats the null* — at round 3 it moves the share $0.11 \rightarrow 0.65$, payoff $0.626 \rightarrow 0.953$ against the null’s 0.706 — and loses on the trace average only because the 16 untouched rounds sit at a near-corner $s = 0.99$ and stay there (payoff 0.199) while the null lifts each of them (to 0.457). Panel B shows this surviving the corrected payoff: even aiming at the wrong point (0.5, not 0.704), the null still gains more. **One trace and two of its rounds, chosen for legibility: an illustration, not a statistical result** — the evidence is the 85-trace paired comparison in this section and Appendix A.

5.2 The reciprocal-framing manipulation: survivorship, and a correction that helps but does not fully resolve it

At 3B, reciprocal framing replicates the pilot’s finding on an independently-drawn budget-set design: GARP pass rate collapses from 0.73 (22/30) to 0.39 (11/28), $p = 0.0089$ (survives Benjamini–Hochberg correction across our full test family, $p_{BH} = 0.011$, but not the more conservative Holm correction, $p_{Holm} = 0.062$); CCEI moves much less and reaches only borderline significance (0.9900 vs. 0.9713, t -test $p = 0.0508$) — the same near-1 compression the pilot flagged, and a further case in this paper of the pattern that GARP pass rate is the sensitive instrument for framing/format disruption while CCEI understates it.

At 1.5B, the manipulation this study was designed around does *not* survive proper correction for discard-selection. First-attempt discard under reciprocal framing at 1.5B was 43.3% (13/30), falling




Figure 3: Top: the 1.5B reciprocal-framing CCEI comparison, pilot (naive single-attempt discard handling, grey) vs. main experiment (capped-retry-corrected, blue). Bottom: the corresponding discard rate at each stage. Neither measurement detects a CCEI shift ($p = 0.66$ pilot, $p = 0.73$ main) — two statistically indistinguishable nulls, not a sign reversal; the pilot’s point estimate reflects which sessions survived, not a real effect that was later corrected away.

to a residual 20.0% (6/30) after the retry protocol — smaller than the pilot’s unretried 52%, plausibly because independent per-replicate budget sets less often reproduce one unusually hard draw, but the pattern is unchanged: reciprocal framing produces far more discards than any other condition, including the *other* 1.5B manipulation (multiturn, 0%). Once the surviving 24 sessions are compared, GARP pass rate (0.40 vs. 0.38, $p = 0.85$) and CCEI (0.9522 vs. 0.9413, $p = 0.73$) are indistinguishable between baseline and reciprocal framing, with confidence intervals overlapping almost entirely. This is no detectable CCEI shift, not a confirmation of the pilot’s own naive point estimate: the pilot’s +0.0169 ($p = 0.66$, itself not significant) and the main experiment’s -0.0109 ($p = 0.73$) are two statistically indistinguishable nulls, and reading the difference in their signs as a reversal of a real effect overstates what either number supports. Figure 3 shows both measurements side by side.

5.3 The multiturn/format effect: a large effect in the opposite direction from the literature

Splitting the 25 rounds into separate sequential calls collapses the GARP pass rate from 40% to 10% ($p = 0.0073$) while CCEI barely moves (0.9522 vs. 0.9540, $p = 0.91$; Figure 4) — a larger drop in percentage points than reciprocal framing produced at 1.5B, with *zero discards on either arm* so no selection-bias caveat applies. Corrected for multiple comparisons across our full test family, this result survives Benjamini–Hochberg correction ($p_{BH} = 0.0094$) but not the more conservative Holm correction ($p_{Holm} = 0.058$); we report both rather than the raw p -value alone. Together with the 3B reciprocal-framing result (§5.2), this is a further case of the same pattern. Against the literature this arm was designed against [Wang et al., 2025a]: in the same budget-allocation domain, on the same model family at a larger scale (Qwen2.5-7B), single-turn elicitation was found to be *more* coherence-degrading than multi-turn; here, at 1.5B, the reverse holds — multi-turn elicitation collapses GARP pass rate relative to our single-turn baseline. We do not have an explanation for the reversal and do not offer one; we report the disagreement as a finding in its own right, not a confirmation of the literature’s own strongest lever.

6 Limitations

CCEI is noisy at 1.5B, and the paper leads with GARP pass rate because of it, not by preference. The 1.5B model’s within-condition CCEI noise (baseline within-subject coefficient of variation 14.8%) is comparable to published human test–retest noise on the same index [Nitsch et al., 2022], and $N = 30$ is explicitly underpowered for a CCEI-scale effect of the pilot’s own magnitude at this scale (the minimum reliably detectable effect at this N is roughly $2.5\times$ the pilot’s magnitude). A CCEI-power-matched design would need $N \approx 111$ to 161 depending on the assumed baseline variance; that scale-up was deferred pending this run’s results and was not executed. We treat this as

1.5B, baseline vs. multiturn format: GARP pass rate and mean CCEI
(zero discards in either arm; error bars are ± 1 SD)




Figure 4: 1.5B, baseline vs. multiturn elicitation format. GARP pass rate (left) collapses; mean CCEI (right, ± 1 SD) is visually flat with heavily overlapping dispersion. Zero discards occurred in either arm, unlike Figure 3.

an open design tradeoff, not a result masked by a metric choice: the pass-rate metric is adequately powered throughout and every CCEI result we report is accompanied by its own confidence interval rather than presented as equally reliable.

Single run per condition per model at $N = 30$. Every (model, condition) cell in this paper is one run of the stated protocol; we do not have repeated full runs of the entire pipeline to report a between-run variance on the headline dose-response statistic itself, only the within-run replicate variance captured by the reported confidence intervals and p -values.

Two model families, one scale point each. The headroom/null-control design isolates the identification strategy from a scale confound but does not trace a dose-response relationship *across* model scale: a pilot found headroom and measurement reliability do not coexist anywhere in the range piloted, so the design targets the scale where headroom exists rather than blending a noisy-but-real candidate effect with a clean-but-absent one across a wider range (a third family was excluded after the multi-model-residency failure noted in §4; extending the roster is a documented follow-up, not a silent omission).

Two fixed exogenous payoffs, not three. We report two independently designed payoffs (§3.3, §3.4) because the first one’s single, globally-identical optimum turned out to be exploitable by a GARP-blind operator; the second removes that specific property and reproduces the same verdict. We do not report a third payoff. A third, run only because the second also came back negative, would be payoff-shopping rather than a robustness check, and §5.1 says so rather than quietly running one until the result flips.

An adverse prior the paper now joins rather than argues against. Three independently published repair attempts moved the wrong way (§2). Our own reciprocal-framing manipulation at 1.5B (§5.2) is a fourth negative in the narrower sense that the manipulation intended to induce measurable coherence variation did not do so once properly measured. The null-operator result of §5.1 is a fifth: under two independently designed exogenous payoffs, restoring GARP-consistency specifically does not outperform a GARP-blind operator matched on displacement. We do not read this as evidence that the coherence-competence relationship is negative in general — C2 remains an open identification question, not a settled sign — only that this paper’s own attempt to find a positive relationship, net of mere displacement, did not succeed.

Broader impacts. This paper’s central result argues directly against imposing GARP-restoration on deployed AI economic agents as a default decision-quality intervention: under two independently designed exogenous payoffs, a GARP-blind operator that merely displaces choices toward an interior point outperforms the real repair (§5.1). A practitioner who deploys coherence-repair expecting a decision-quality gain, on the strength of a raw dose-response correlation alone, would be adopting exactly the confound this paper’s control was built to catch. The reciprocal-framing manipulation this study was designed around similarly turned out, once corrected, to have no measurable effect on the coherence it was thought to disrupt (§5.2). The paper’s contribution is a measurement design for

asking whether repair helps in a given setting — and, in the setting it tests, a documented instance where it does not — not a recommendation to deploy it universally, and not a demonstration that no coherence-repair method could ever help downstream decisions. Separately, the discard-selection finding (§4.3) generalizes beyond this paper: any revealed-preference or consistency instrument that silently drops unparseable or disruption-caused failures risks systematically underestimating exactly the manipulations that matter most, which bears on how AI economic-behavior audits are designed and reported more broadly. We see no elevated misuse risk specific to this work: it uses only synthetic budget-allocation tasks on publicly available open-weight models, releases no new model or dataset, and does not involve human subjects.

7 Conclusion

This paper set out to measure whether restoring GARP-consistency in an LLM agent’s own realized choices buys more exogenous payoff than the mechanical fact of displacement alone. Applying Demuynck and Rehbeck [2023]’s minimal-quantity-error GARP repair post hoc to a frozen agent, we find a real, precisely estimated, monotone raw dose–response relationship between repair size and payoff gain. But a distance-matched, GARP-blind null operator — absent from every prior axiom-enforcement result we are aware of — outperforms the real repair on that same payoff, significantly, at both model scales; a second, independently designed payoff built specifically to remove the first one’s exploitable structure reproduces the same verdict. C1, as originally proposed, is not supported: the answer this paper’s control was built to give is no. A partial mechanistic finding — trace extremity predicts the null’s advantage in both experiments — offers a candidate explanation without fully resolving it. What survives is the paper’s identification argument itself (C2): the coherence–competence relationship remains open, not settled in either direction, and this is the first attempt we know of to measure it without confounding coherence with capacity, a preference-judgment outcome, or mere displacement magnitude. Two further, unpredicted findings — a purpose-built framing manipulation’s apparent 1.5B effect was a discard-selection artifact rather than a real coherence shift (C3), and an elicitation-format manipulation produces a large GARP-pass-rate collapse in the opposite direction from an independently published finding in the same domain and model family at a larger scale — are reported in their own right. A well-controlled negative, an open identification question, and an instrument-validity finding are exactly the kind of result this literature should surface by default rather than paper over.

References

- S. N. Afriat. The construction of utility functions from expenditure data. *International Economic Review*, 8(1):67–77, 1967.
- Victor H. Aguiar and Nail Kashaev. GARP-EFM: Improving foundation models with revealed preference structure. *arXiv preprint arXiv:2603.23993*, 2026.
- James Andreoni and John Miller. Giving according to GARP: An experimental test of the consistency of preferences for altruism. *Econometrica*, 70(2):737–753, 2002.
- Isaiah Andrews. Revealed rationality: Label-free evaluation and regularization from representation theorems. *arXiv preprint arXiv:2608.05015*, 2026.
- Stephen G. Bronars. The power of nonparametric tests of preference maximization. *Econometrica*, 55(3):693–698, 1987.
- Joy Buchanan and Joshua Foster. The innate economic preferences of language models. *arXiv preprint arXiv:2607.26288*, 2026.
- Alina Chadwick, Anson Kahng, and Jens Kipper. Dutch books and money pumps: rectifying vulnerabilities in LLMs through rationality. In *Proceedings of the 4th International Conference on Human and Artificial Rationality (HAR)*, 2025.
- Yiting Chen, Tracy Xiao Liu, You Shan, and Songfa Zhong. The emergence of economic rationality of GPT. *Proceedings of the National Academy of Sciences*, 120(51):e2316205120, 2023.

493 Syngjoo Choi, Raymond Fisman, Douglas Gale, and Shachar Kariv. Consistency and heterogeneity
494 of individual behavior under uncertainty. *American Economic Review*, 97(5):1921–1938, 2007.

495 Syngjoo Choi, Shachar Kariv, Wieland Müller, and Dan Silverman. Who is (more) rational? *American*
496 *Economic Review*, 104(6):1518–1550, 2014.

497 Thomas R. Cook, Sophia Kazinnik, Zach Modig, and Nathan M. Palmer. What do LLMs want?
498 Technical Report FEDS 2026-006, Federal Reserve Board, 2026.

499 Mark Dean and Daniel Martin. Measuring rationality with the minimum cost of revealed preference
500 violations. *Review of Economics and Statistics*, 98(3):524–534, 2016.

501 Thomas Demuynck and John Rehbeck. Computing revealed preference goodness-of-fit measures
502 with integer programming. *Economic Theory*, 76(4):1175–1195, 2023.

503 M. Houtman and J. A. Maks. Determining all maximal data subsets consistent with revealed
504 preference. *Kwantitatieve Methoden*, 19:89–104, 1985.

505 Huang, Li, Zhao, and Li. Transitivity meets cyclicity: Explicit preference decomposition for dynamic
506 llm alignment. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2605.17342.

507 Q. Huangfu and J. A. J. Hall. Parallelizing the dual revised simplex method. *Mathematical Program-*
508 *ming Computation*, 10(1):119–142, 2018.

509 Felix J. Nitsch, Luca M. Lüpken, Nils Lüscho, and Tobias Kalenscher. On the reliability of
510 individual economic rationality measurements. *Proceedings of the National Academy of Sciences*,
511 119(31):e2202070119, 2022.

512 Shumiao Ouyang, Hayong Yun, and Xingjian Zheng. AI as decision-maker: Ethics and risk prefer-
513 ences of LLMs. *arXiv preprint arXiv:2406.01168*, 2024.

514 Mingxuan Sun. CONSISTRE: A unified consistency-aware framework for document-level relation
515 extraction with large language models. *arXiv preprint arXiv:2607.24312*, 2026.

516 Hal R. Varian. The nonparametric approach to demand analysis. *Econometrica*, 50(4):945–973,
517 1982.

518 Wang, Yao, Zhang, Gai, Liu, and Zhong. When experimental economics meets large language
519 models: Tactics with evidence. *arXiv preprint arXiv:2505.21371*, 2025a.

520 Wang, Zhan, et al. Trustroboreward: Preference-ordered isotonic score editing for multi-paradigm
521 robot reward modeling. *arXiv preprint arXiv:2608.08491*, 2026.

522 Yidong Wang, Yunze Song, Tingyuan Zhu, Xuanwang Zhang, Zhuohao Yu, Hao Chen, Chiyu
523 Song, Qiufeng Wang, Cunxiang Wang, Zhen Wu, Xinyu Dai, Yue Zhang, Wei Ye, and Shikun
524 Zhang. Trustjudge: Inconsistencies of LLM-as-a-judge and how to alleviate them. *arXiv preprint*
525 *arXiv:2509.21117*, 2025b.

526 Stephen Warshall. A theorem on Boolean matrices. *Journal of the ACM*, 9(1):11–12, 1962.

527 Kaiser Yamin, J. Tang, S. Cortes-Gomez, A. Sharma, Eric Horvitz, and Bryan Wilder. When
528 agents say one thing and do another: Validating elicited beliefs from LLMs. *arXiv preprint*
529 *arXiv:2602.06286*, 2026a.

530 Khurram Yamin, Jingjing Tang, Eric Horvitz, and Bryan Wilder. Can revealed preferences clarify
531 LLM alignment and steering? *arXiv preprint arXiv:2605.08556*, 2026b.

532 Yifan Zeng, Ojas Tendolkar, Raymond Baartmans, Qingyun Wu, Lizhong Chen, and Huazheng
533 Wang. LLM-RankFusion: Mitigating intrinsic inconsistency in LLM-based ranking. *arXiv preprint*
534 *arXiv:2406.00231*, 2024.

535 Zhang, Zhang, Wu, Xu, and Gu. Beyond bradley-terry models: A general preference model for
536 language model alignment. In *International Conference on Machine Learning (ICML)*, 2025.
537 arXiv:2410.02197.

538 Jian-Qiao Zhu, Haijiang Yan, and Thomas L. Griffiths. Recovering event probabilities from large
539 language model embeddings via axiomatic constraints. *arXiv preprint arXiv:2505.07883*, 2025.

540 A Payoff exogeneity audit

541 Full detail of the four-part adversarial audit summarized in §3.3.

542 **(1) No shared derived quantity beyond the exogenous (p, I) .** The payoff implementation has no
543 dependency on the projection implementation; the optimal bundle x_t^* is a function of (p_t, I_t) alone
544 and never references x_t or $\tilde{x}_t$. The one near-miss is that the projection’s feasibility incumbent is a
545 Cobb–Douglas demand *share-fitted to the agent’s own observed data* — a plausible leakage channel
546 on first read — but this quantity is never passed to the underlying MILP solver as an actual incumbent
547 (the solver used has no such interface); it is used only to log an upper-bound sanity ceiling on the
548 reported distance, never to steer the solution.

549 **(2) The payoff’s optimum is independent of revealed preference, checked empirically as well
550 as by construction.** Comparing raw (pre-repair) payoff between the 57 traces that were already
551 GARP-consistent and the 85 that were not: mean 0.7872 vs. 0.7899, Welch $t = -0.07$, $p = 0.94$.
552 GARP-consistency status alone does not predict payoff, ruling out a population-level mechanical link
553 between coherence and payoff score.

554 **(3) Hand-derived mechanism check.** Five traces spanning the full dose range (0.17 to 111.64)
555 were inspected directly, comparing the direction of the actual projection movement $(\tilde{x} - x)$ against
556 the direction toward the exogenous optimum $(x^* - x)$ via cosine similarity. Similarities were low
557 (0.16–0.31) on the four traces where payoff improved after repair, and *negative* (−0.35) on the single
558 trace where payoff worsened — the sign of the outcome tracks a geometric quantity the projection
559 objective never references, evidence against the mechanism being a disguised optimization toward
560 the payoff target.

561 **(4) The severity-confound check.** Larger-dose traces do start from worse raw payoff (Pearson
562 $r = -0.41$, $p < 0.0001$), and worse-starting traces do have more room to improve ($r = -0.41$
563 between raw payoff and Δ payoff) — a genuine ceiling-effect confound. It does not explain the
564 dose–response relationship away: the partial correlation of dose against Δ payoff controlling for raw
565 payoff is $r = 0.784$ ($p = 7 \times 10^{-19}$, attenuated only slightly from the unconditional $r = 0.821$), and
566 in a joint OLS regression the dose coefficient is highly significant ($t = 11.43$, $p < 10^{-16}$) while raw
567 payoff’s own coefficient is not ($t = -1.36$, $p = 0.18$).

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and §1 state exactly what is claimed (a measured dose–response relationship, an instrument-validity finding, and a replication of an independently published format effect) and what is not (novelty of repair, of projection, or of the sign of the effect); §6 states the design’s scope limits explicitly.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: §6 (Limitations) discusses CCEI’s noise and power at 1.5B, the single-run-per-condition design, the two-model-family scope, the single fixed payoff weighting, and the adverse prior from three prior negative results.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper's contribution is a computational method (a MILP formulation for GARP-repair) building on a cited external representation theorem [Demuyne and Rehbeck, 2023]; it does not state or prove new theorems of its own.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §4-5 specify the exact models, budget-set generator and seeding scheme, conditions, power calculation, retry protocol, and the full MILP formulation (§3) needed to reproduce the main results; the exogenous payoff is a closed-form function stated in full.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

673 (d) We recognize that reproducibility may be tricky in some cases, in which case
674 authors are welcome to describe the particular way they provide for reproducibility.
675 In the case of closed-source models, it may be that access to the model is limited in
676 some way (e.g., to registered users), but it should be possible for other researchers
677 to have some path to reproducing or verifying the results.

678 5. Open access to data and code

679 Question: Does the paper provide open access to the data and code, with sufficient instruc-
680 tions to faithfully reproduce the main experimental results, as described in supplemental
681 material?

682 Answer: [Yes]

683 Justification: An anonymized code and data package is included as supplementary material
684 (prompts, seeds, the parser/validity rule, and raw per-trace logs for the main experiment, the
685 null-operator control, the corrected-payoff robustness check, and the pilot), with a README
686 giving the exact commands to regenerate every reported statistic from the included raw logs
687 without re-running any model inference. Both LLMs used are publicly available open-weight
688 models, so the full pipeline is also reproducible from scratch.

689 Guidelines:

- 690 • The answer [N/A] means that paper does not include experiments requiring code.
- 691 • Please see the NeurIPS code and data submission guidelines ([https://neurips.cc/](https://neurips.cc/public/guides/CodeSubmissionPolicy)
692 [public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 693 • While we encourage the release of code and data, we understand that this might not
694 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
695 including code, unless this is central to the contribution (e.g., for a new open-source
696 benchmark).
- 697 • The instructions should contain the exact command and environment needed to run to
698 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 699 • The authors should provide instructions on data access and preparation, including how
700 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 701 • The authors should provide scripts to reproduce all experimental results for the new
702 proposed method and baselines. If only a subset of experiments are reproducible, they
703 should state which ones are omitted from the script and why.
- 704 • At submission time, to preserve anonymity, the authors should release anonymized
705 versions (if applicable).
- 706 • Providing as much information as possible in supplemental material (appended to the
707 paper) is recommended, but including URLs to data and code is permitted.

709 6. Experimental setting/details

710 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
711 rameters, how they were chosen, type of optimizer) necessary to understand the results?

712 Answer: [Yes]

713 Justification: §4 specifies model identities, all three conditions' exact framing, the budget-set
714 generator and its rejection rule, seed scheme, replicate count and power calculation, and the
715 discard/retry protocol.

716 Guidelines:

- 717 • The answer [N/A] means that the paper does not include experiments.
- 718 • The experimental setting should be presented in the core of the paper to a level of detail
719 that is necessary to appreciate the results and make sense of them.
- 720 • The full details can be provided either with the code, in appendix, or as supplemental
721 material.

722 7. Experiment statistical significance

723 Question: Does the paper report error bars suitably and correctly defined or other appropriate
724 information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Every reported comparison in §5 carries a p -value, and Table 1 and the dose-response statistics report means alongside sample sizes; CCEI comparisons additionally report standard deviations and, in §6, confidence-interval-based power limitations.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All inference was run locally on a single consumer CPU-class machine (no GPU, no commercial API calls), with the full $N = 30$ design measured at approximately 2.1 hours wall-clock (§4.1); every MILP projection solves in under 5 seconds with MIP gap $\leq 8.1 \times 10^{-5}$.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The study elicits only synthetic, non-sensitive budget-allocation choices from open-weight language models; it involves no human subjects, no personal data, and no scraped content.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: See the Broader impacts paragraph preceding §7: the paper’s own finding that 18% of repairs worsen the exogenous payoff is itself a caution against treating coherence-repair as a costless, universally beneficial deployment default.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no model or dataset; both LLMs used are existing, publicly released open-weight models used unmodified as measurement subjects.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The two LLMs used are named with version in §4.1. Qwen2.5-1.5B-Instruct (Alibaba Cloud) is released under the Apache License 2.0. Llama-3.2-3B-Instruct (Meta Platforms, Inc.) is released under the Llama 3.2 Community License Agreement, a custom commercial license. Both are publicly available open-weight models used here for research purposes consistent with their license terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: This submission does not release a new dataset or model asset; the generated budget-set choice logs are experimental records, not a packaged new asset.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The study involves no crowdsourcing and no human subjects; all choices are elicited from language models.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The study involves no human subjects and required no IRB review.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: The two LLM agents under study are the paper's core experimental subject, not a writing aid: §4.1 names both models and their role as the choice-making agents whose revealed preferences are elicited, projected, and scored.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.