
What Does Repairing Choice Inconsistency Actually Buy?

A Budget-Set Diagnosis

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Repairing an AI agent’s incoherent preferences is not new: inference-time transi-
2 tivity repairs, isotonic projections onto preference-defined cones, and training-time
3 rationality penalties have all been proposed, several with reported downstream
4 gains. None controls for displacement magnitude: no result establishes whether
5 restoring coherence itself, not mere displacement toward an interior point, is what
6 any reported gain comes from. We build that control. Given an LLM agent’s own
7 choices over budget sets, we compute the minimal MILP perturbation restoring
8 GARP-consistency, then compare it against a GARP-blind null operator spending
9 the identical displacement budget shrinking every bundle toward a fixed center.
10 Under the paper’s original exogenous payoff, the null outperforms the real repair by
11 more than 2x (mean payoff gain 0.0220 vs. 0.0091, Wilcoxon $p = 3.98 \times 10^{-10}$,
12 winning 81% of traces). A second, independently designed payoff built to re-
13 move the first one’s exploitable fixed target reproduces the same verdict (null
14 0.0217 vs. real 0.0072, $p = 1.24 \times 10^{-5}$, 70.8% win rate); we decline a third as
15 payoff-shopping. Trace extremity predicts the null’s advantage in both experiments.
16 The paper’s standing contributions are instead a capacity-deconfounded identifica-
17 tion strategy for the coherence–competence question, which remains open, and a
18 discard-selection instrument-validity finding.

19 1 Introduction

20 An LLM agent asked to allocate a fixed budget across two goods at varying prices, repeatedly, will
21 frequently violate the Generalized Axiom of Revealed Preference (GARP; Afriat, 1967, Varian, 1982):
22 it prefers bundle x_1 to x_2 at one budget line and the reverse at another, in a way no monotone utility
23 function can rationalize. The same budget-allocation instrument is the standard paradigm for testing
24 GARP experimentally on human subjects [Andreoni and Miller, 2002, Choi et al., 2007, 2014], and a
25 growing literature applies it to GPT-family models directly, without a repair step, to ask whether they
26 satisfy classical rationality axioms at all [Chen et al., 2023]. A separate, growing literature repairs
27 exactly this kind of inconsistency in LLM systems — restoring transitivity in pairwise judgments,
28 projecting cardinal scores onto monotone cones, penalizing incoherence during training — and
29 several of these systems report that the repaired outputs are downstream-better than the raw ones.
30 What none of them establishes is a *dose*: how much better, as a function of how much repair was
31 needed, isolated from the capacity and training confounds that come from comparing a richer model
32 class to a poorer one.

33 We study this question in a setting where it can be answered cleanly. An LLM agent makes a
34 sequence of budget-constrained choices over two goods. We compute, as a single mixed-integer
35 linear program, the minimal quantity perturbation that would make the agent’s own realized choice

sequence GARP-consistent — a projection applied *post hoc* to a frozen agent, so nothing about its weights, training, or decoding changes across doses. The size of that perturbation is a graded, cardinal dose. We score both the raw and the repaired sequence against a fixed, equal-weight Cobb–Douglas payoff whose optimum at each budget line is a closed-form function of prices and income alone — never of the agent’s own choices — so the outcome measure is exogenous to the preference data by construction, not a restatement of coherence. Tracing dose against Δ payoff gives the relationship this paper reports.

That relationship turns out not to isolate what it was built to isolate. A control absent from every published axiom-enforcement result we are aware of — an operator that knows nothing about GARP but spends the identical displacement budget shrinking every bundle toward a fixed interior point — outperforms the real GARP-restoring repair on the exogenous payoff, significantly, at both model scales, under two independently designed payoffs (§5.1). This is the paper’s central finding: restoring rationality, as originally proposed, does not buy more exogenous payoff than mere displacement toward the interior of the budget line does on its own. We report this as a negative result reached twice by construction, not as a null we happened to observe once.

We do not claim any of the individual pieces are new; §2 concedes exactly what is not. What has not been done is the conjunction: an agent’s own realized choices, a graded coherence-indexed dose, and an outcome that is not itself a preference judgment, run together so that the identification is clean. We run this at a model scale where a pilot study found measurable coherence headroom (1.5B parameters), and repeat the identical pipeline at a scale with almost none (3B parameters) as an identification control on the method itself, not as a second point on a shared dose curve across scale.

Two further results were not predicted by the design. The reciprocal-price framing manipulation intended to induce coherence variation at 1.5B produced, in a pilot run, a large apparent effect that disappears once discard-selection is corrected with a capped retry protocol: the pilot’s naive estimate and the corrected main-experiment estimate are both statistically indistinguishable from zero. We report the discard rate itself, not the CCEI point estimate, as the more reliable signal that the manipulation disrupted the agent (our claim C3), quantifying how much a naive discard rule can distort a measured effect (§5.2). Separately, splitting single-turn elicitation into separate sequential calls at 1.5B produces a large GARP-pass-rate collapse, moving in the *opposite* direction from an independently published finding in the same domain and model family at a larger scale, with no discard confound on either arm.

Contributions. (1) An application — not a novel method — of Demuynck and Rehbeck [2023]’s minimal-quantity-error MILP projection to an agent’s own realized choices, independently verified on every output rather than trusted from solver status alone. (2) A closed-form exogenous payoff, fixed before any data collection, audited adversarially for leakage and for the mechanical confound that larger violations simply have more room to improve. (3) A distance-matched, GARP-blind null-operator control — absent from every published axiom-enforcement result we are aware of — run against two independently designed exogenous payoffs, showing that the measured raw dose–response relationship is not evidence that GARP-restoration specifically, rather than mere displacement toward the interior, is what buys the exogenous payoff gain (§5.1). (4) An instrument-validity finding (discard-selection bias under a disruptive manipulation), corrected for by a stated retry protocol and reported per cell rather than pooled away. (5) A single-turn-vs-multi-turn elicitation-format manipulation that produces a large GARP-pass-rate collapse at 1.5B, in the *opposite* direction from an independently published single-turn-vs-multi-turn finding in the same budget-allocation domain and the same model family (Qwen2.5) at a larger scale.

2 Related Work

Repairing an AI system’s incoherent preferences is not new; Table 1 positions nine published systems against the four criteria that jointly define our claim: an agent’s own choices (not a judgment about third-party items), an exogenous payoff (no preference judgment enters it), a graded rather than binary dose, and a formal minimum-distance projection. None occupies all four.

POISE [Wang et al., 2026] is the sharpest vocabulary collision: pool-adjacent-violators projection onto a closed convex chain-monotone cone, with a proved guarantee the edit lands weakly closer to a posited ground truth. We concede the minimum-distance priority unqualified: projection onto a closed convex set is non-expansive in L_2 ; the GARP-consistent set is a union of polyhedra, hence not convex,

Table 1: Where prior repair systems sit. *Own choices*: the repaired object is a choice the agent made for itself from a budget set, not a judgment about third-party items. *Exogenous payoff*: an outcome into which no preference judgment enters.

	Own choices	Exogenous payoff	Graded dose	Min.-distance projection
POISE [Wang et al., 2026]	no	no	no	yes
Rationality layer [Chadwick et al., 2025]	no	no	no	yes
TrustJudge [Wang et al., 2025b]	no	no	no	no
CONSISTRE [Sun, 2026]	no	partial	no	no
LLM-RankFusion [Zeng et al., 2024]	no	yes	no	no (declined)
Innate preferences [Buchanan and Foster, 2026]	yes	no	no	no
GARP-EFM [Aguilar and Kashaev, 2026]	no	yes	no	no
HRC/DSPPO [Huang et al., 2026]	no	no	yes	no
Control vectors [Cook et al., 2026]	yes	no	yes	no
This paper	yes	yes	yes	yes

so no analogue transfers (§B). POISE also projects a *teacher’s* offline labels, not an agent’s own choices. Three further inference-time repairs lack a genuine degree parameter or Afriat machinery, on different objects: Chadwick et al. [2025], TrustJudge [Wang et al., 2025b], CONSISTRE [Sun, 2026]. LLM-RankFusion [Zeng et al., 2024] has a genuinely exogenous payoff but declines a minimum-distance method. Two training-time systems, one altering the agent itself [Buchanan and Foster, 2026] and one not an LLM agent at all [Aguilar and Kashaev, 2026], report no downstream comparison and a forecasting evaluation respectively.

A parallel line varies the transitivity of the *preference model* rather than an agent’s choices: GPM [Zhang et al., 2025] and HRC/DSPPO [Huang et al., 2026] both make the cycle-tolerant arm a strict superset model class, confounding coherence with capacity by construction — and on GPM’s own length-controlled metric that arm *loses* in 18 of 24 cells. HRC/DSPPO’s own inverted-U dose-response curve is over a training-time schedule weight on a learned proxy, scored by an LLM judge, where ours acts post hoc and is exogenous. Our closest theoretical neighbour, Andrews [2026], proposes 1 – CCEI as a training penalty argued a priori, never measured; ours is the empirical counterpart to a proposal that has circulated unrun.

The sharpest objection is psychometric, and it is in PNAS. Nitsch et al. [2022] find CCEI/Houtman-Maks [Houtman and Maks, 1985] reliability never reaches acceptable levels across eight datasets, and that participants given the chance to revise their own choices saw mean CCEI *fall*. That revision arm is not a repair operator (a random subset, no consistency objective, no guarantee of fewer violations) — ours attains rationalizability by construction, the same distinction that disposes of Yamin et al. [2026]’s isotonic repair (worse in 14 of 16 cells). Nitsch et al. diagnose their finding as a between-subject-variance problem whose prescribed remedy — a between-groups design — is this paper’s design; three independently published axiom-enforcement results (adding Zhu et al. [2025]’s) all point the same adverse way, which is why we carry a null-operator control neither published negative had.

The closest neighbour on the economics side is invisible to any arXiv sweep. Cook et al. [2026] steer economic choices toward payoff-maximising behaviour via personas and control vectors — occupying two of our three legs without Afriat machinery; their finding that reframing moves models while persona prompting does not corroborates our manipulation choice, anchored in Wang et al. [2025a], who find *format* moves the Afriat index in the same domain and model family (Qwen2.5) at a larger scale: collapsing multi-turn to single-turn drops CCEI by up to 0.241. Our opposite manipulation finds the opposite direction (GARP pass rate $0.40 \rightarrow 0.10$, $p = 0.0073$, zero discards) — not a replication, since which format is more degrading reverses between their 7B model and ours.

3 Method

3.1 The projection: Demuynck & Rehbeck’s (2023) minimal-quantity-error MILP, applied and independently verified

This section applies, rather than proposes, a projection method. The MILP formulation below is Demuynck and Rehbeck [2023]’s multiplier-free ordinal characterization of minimal-quantity-error GARP repair, adopted here on an LLM agent’s own choice sequence and verified independently on every output; nothing about the formulation itself is new to this paper.

For a sequence of T observations $(p_t, x_t)_{t=1}^T$, K goods, GARP holds if there is no sequence of choices such that x_{t_1} is revealed preferred to x_{t_2} , ..., x_{t_n} is revealed preferred to x_{t_1} , with at least one strict revealed preference in the cycle. Given a GARP-violating sequence, we seek the minimal perturbation $\tilde{x}_t$ that restores GARP-consistency. The naive Afriat-multiplier formulation is bilinear in the unknowns once $\tilde{x}$ is itself a decision variable; we instead use the multiplier-free ordinal characterization of Demuynck and Rehbeck [2023], under which every constraint is linear. At $T = 25$ this is 600 binaries and 125 continuous variables per trace, solved as a single MILP on the HiGHS backend [Huangfu and Hall, 2018]. The full formulation, and the three method commitments governing budget exhaustion, the strict-preference margin, and per-output verification, are given in Appendix B; every one of the 85 GARP-violating traces’ projections was independently re-verified GARP-consistent via the same combinatorial check used throughout (MIP gap $\leq 8.1 \times 10^{-5}$ on every solve).

The objective is L_1 : minimal total absolute displacement between $\tilde{x}$ and x across goods and rounds, with uniform weights, keeping the program a pure MILP (L_∞ is recorded alongside but not separately analyzed; L_2 would require an MIQP solver we do not have configured; the exact objective and constraint definitions are in Appendix B). The *dose* for a trace is its L_1 projection distance $\sum_t \sum_k |x_{t,k} - \tilde{x}_{t,k}|$, comparable across sessions without further normalization because income is fixed throughout — the same style of graded, minimum-cost severity measure as Dean and Martin [2016], applied here to an LLM’s own choices and paired with an exogenous behavioral payoff rather than left as a standalone index. A Cobb–Douglas demand share-fitted to each trace’s own observed data is computed as a feasibility incumbent and sanity ceiling on the reported distance, recorded only as an upper-bound diagnostic; Appendix A(1) confirms it never reaches the solver.

What can be guaranteed, and what cannot, relative to the closest published structural analogue [Wang et al., 2026], is addressed in Appendix B: because the GARP-consistent set is a union of polyhedra rather than a single convex set, no distance-minimization guarantee analogous to a convex-cone projection is claimed.

3.2 The exogenous payoff

The payoff must not be derived from the agent’s own revealed choices — a utility function fit to the agent’s choices (such as the projection’s own feasibility incumbent above) would be circular if used as the outcome measure, rewarding the agent for resembling its own revealed preferences rather than measuring anything external to them. We instead fix, before any model was queried and never re-estimated from any agent’s choices, an equal-weight Cobb–Douglas valuation $U_{\text{exo}}(x) = x_A^{0.5} x_B^{0.5}$. At budget line (p_t, I_t) the optimal bundle x_t^* is the closed-form Cobb–Douglas demand $x_{t,A}^* = 0.5I_t/p_{t,A}$, $x_{t,B}^* = 0.5I_t/p_{t,B}$ — a function of prices and income alone, never of the agent’s observed or projected bundle. The payoff score is the efficiency ratio $\text{payoff}(x_t) = U_{\text{exo}}(x_t)/U_{\text{exo}}(x_t^*) \in (0, 1]$, computed identically for the raw and projected sequence at the same budget lines, so that $\Delta \text{payoff} = \text{payoff}(\tilde{x}) - \text{payoff}(x)$ is comparable across conditions, models, and repairs. This mirrors — not the specifics, but the structure of — a money-metric utility index, the standard welfare-comparison device over the same GARP/Afriat objects this paper’s method already relies on. A genuinely exogenous, real-world payoff of this kind is rare in this literature: Ouyang et al. [2024] score LLM-generated forecasts against firms’ realized future spending and find a *non-monotonic* alignment-accuracy relationship, but without a placebo of matched size cannot separate an alignment-specific effect from a generic fine-tuning effect — the same separation our design is built to make.

We audited this design adversarially, before treating any dose–response result as reportable, for four failure modes: a shared-data leak between payoff and projection, a mechanical link between GARP-consistency and payoff, a projection direction secretly aimed at the payoff optimum, and the




Figure 1: The pipeline, one replicate slot end to end: budget-set draw (1); prompt/format arm (2); parse under the capped retry protocol of §4.3, residual discards excluded (3); combinatorial Warshall-closure GARP check (4); and, for the traces that violate it, the minimal-quantity-error MILP of §3.1 (5), whose L_1 displacement is the dose. Steps 6–7: the real repaired sequence (6a, solid blue) and the GARP-blind null spending an identical L_1 displacement (6b, dashed orange) are scored (7) under both payoff designs. Counts are the main experiment’s (§5).

176 confound that larger-dose traces simply start worse and have more room to improve. None survived
 177 the audit (Appendix A gives every check and number); the raw dose–response statistic we report in
 178 §5.1 is not an artifact of any of the four. A fifth failure mode — that this payoff’s own concavity
 179 could reward mere displacement without any GARP-specific effect — required a new control rather
 180 than an audit of the existing design, and is addressed separately in §3.3.

181 3.3 Two robustness controls against a payoff-geometry confound

182 The payoff of §3.2 reduces to $\text{payoff}(s) = 2\sqrt{s(1-s)}$, s the expenditure share on good A: concave,
 183 single interior maximum at $s = 0.5$, identical across every trace. Any operator that moves a bundle
 184 toward that fixed point raises its payoff regardless of GARP. We construct a **null operator**: for
 185 each GARP-violating trace, shrink every observed bundle toward the exogenous optimum x_t^* by
 186 a single shared factor $\lambda \in [0, 1]$ chosen so the null’s total L_1 displacement exactly matches the
 187 real projection’s own dose. The null never calls the GARP check and does not, in general, restore
 188 consistency; it is scored by the same payoff function used everywhere else.

189 Because the exploitable property is a single fixed target identical across every trace, we build a
 190 second, independently designed payoff that removes it: a per-trace Cobb–Douglas weight $\alpha_s \sim$
 191 $\text{Uniform}(0.05, 0.95)$, drawn independently for each of the 85 traces and each of $K = 20$ replicate
 192 draws, read only at scoring time. We report two null constructions under this payoff, never pooled: a
 193 **primary**, information-fair null (unaware of α_s) and an **oracle** null (aims at the true per-trace optimum
 194 $x_{\alpha_s}^*$), reported only as an upper bound since it uses information no real repair algorithm has. Figure 1
 195 summarizes the full pipeline.

196 4 Experimental design

197 4.1 Models and conditions

198 Two locally-hosted open-weight models, run entirely on local compute at zero API cost:
199 `qwen2.5:1.5b-instruct` (the headroom model, per a pilot with measurable within-condition
200 coherence variation) and `llama3.2:3b-instruct` (the null-effect control, per the same pilot finding
201 almost no headroom — mean baseline CCEI 0.99). A third model family was excluded after a
202 reproducible multi-model-residency crash under sustained rotation; the two-model, model-major
203 design is both a wall-clock optimization and, after that failure, a correctness requirement. **Compute.**
204 All inference ran locally via 4-bit-quantized weights, no GPU or commercial API calls; the full
205 $N = 30$ design completed in ≈ 2.1 hours wall-clock, every MILP projection (§3.1) solving in under
206 5 seconds.

207 Three conditions at 1.5B, two at 3B. **Baseline:** direct per-token-return framing, single turn, all 25
208 rounds in one prompt and response. **Reciprocal:** identical single-turn format with an inverted-
209 price framing manipulation, the strongest manipulation identified in the pilot. **Multiturn** (1.5B
210 only): baseline framing, but each of the 25 rounds is a separate sequential call carrying the running
211 conversation history, isolating the elicitation-format mechanism from the price-framing mechanism;
212 motivated by an independently published finding, in the same domain, that collapsing multi-turn
213 to single-turn moves CCEI by up to -0.241 for a same-family model at a larger scale [Wang et al.,
214 2025a]. Our arm tests the opposite direction (single-turn split into multi-turn), the direction our own
215 baseline already uses. Not added at 3B, whose role is a null-effect control rather than a second scale
216 point.

217 4.2 Budget sets and power

218 Each of $N = 30$ independent replicates per (model, condition) cell draws its own fresh $T =$
219 25 , $K = 2$ budget set (continuous-density prices, $M, N \sim U[0.1, 1.0]$ i.i.d. rejected unless
220 $\max(M, N) \geq 0.5$, budget exhaustion enforced), seeded independently — unlike a pilot’s single
221 shared set, needed because the main experiment targets independent between-condition replication.
222 Continuous, independently-drawn prices with enforced budget exhaustion also avoid a known pathol-
223 ogy: CCEI can read exactly 1.0 on GARP-violating data whenever a comparison lands on exact
224 budget-line equality, common under grid-sampled or rounded prices [Andrews, 2026]. $N = 30$ gives
225 power ≥ 0.999 against the pilot’s full framing-effect magnitude on the pass-rate metric and ≥ 0.80
226 down to $\sim 60\%$ attenuation, but is explicitly underpowered for a CCEI-scale effect of the pilot’s
227 magnitude at 1.5B (minimum detectable ≈ 0.11), reported rather than papered over. Every replicate’s
228 Bronars power [Bronars, 1987] is recorded; all 150 draws held power ≥ 0.998 .

229 4.3 The discard-selection problem, and its correction as a stated contribution

230 A pilot run found that under reciprocal framing at 1.5B, 52% of sessions (13 of 25) failed the
231 required output-format contract and were silently discarded under the pilot’s own naive handling
232 — standard practice, and one this paper measures the consequences of directly. This is a claim
233 about instrument validity in its own right, true regardless of whether projection helps, hurts, or does
234 nothing downstream. The main experiment implements a capped retry protocol: a session failing
235 the $\geq 20/25$ -valid-rounds threshold is retried up to three total attempts with a fresh seed offset; a
236 slot still failing after three attempts is a residual discard, excluded from CCEI/GARP but retained for
237 audit.

238 Table 2 breaks every reciprocal-framing session down by attempt outcome — kept on the first
239 attempt, rescued by a later retry, or still discarded after three attempts — for both models. At 1.5B,
240 retry-rescued sessions are not a close stand-in for first-attempt sessions: higher mean CCEI but a
241 substantially lower GARP pass rate, and the residual-discard group scores lowest on the handful that
242 can be measured — consistent with, though not conclusive given the small residual sample, some
243 of the same selection concern the pilot’s naive handling raised persisting at a smaller scale even
244 after correction. At 3B, the first-attempt and retry-rescued groups sit much closer together, and the
245 selection concern shows up more sharply at 1.5B than at 3B. We report first-attempt and residual
246 post-retry discard rates as a first-class per-condition outcome for every cell.

Table 2: Discard-outcome breakdown for all 60 reciprocal-framing sessions (30 per model), by attempt group: kept on the first attempt, rescued by a later retry, or still discarded after three attempts (§4.3). GARP pass and CCEI use the same methodology as Table 3; mean dose (L_1) is averaged over the full group, with 0 for any already-GARP-consistent trace. *qwen residual discard: 3 of the 6 slots returned zero valid rounds on every attempt (no p/x data at all); the reported CCEI, GARP pass rate, and dose for this row are computed over the remaining 3 slots only.

Model	Attempt group	n	mean CCEI	GARP pass	mean dose (L_1)
llama3.2:3b	first-attempt success	21	0.9742	0.4286	5.086
llama3.2:3b	retry-rescued	7	0.9624	0.2857	9.839
llama3.2:3b	residual discard	2	1.0000	1.0000	0.000
qwen2.5:1.5b	first-attempt success	17	0.9315	0.4706	13.847
qwen2.5:1.5b	retry-rescued	7	0.9651	0.1429	7.600
qwen2.5:1.5b	residual discard	6	0.8821*	0.6667*	24.196*

Table 3: Headline results per (model, condition) cell. GARP pass and CCEI are computed on kept traces only; dose and Δ payoff are computed on the subset of kept traces that violate GARP.

Model	Condition	n kept	mean CCEI	GARP pass	mean dose (L_1)	mean Δ payoff
llama3.2:3b	baseline	30	0.9900	0.73	5.01	+0.0018
llama3.2:3b	reciprocal	28	0.9713	0.39	6.27	+0.0016
qwen2.5:1.5b	baseline	30	0.9522	0.40	16.68	+0.0090
qwen2.5:1.5b	multiturn	30	0.9540	0.10	14.54	+0.0083
qwen2.5:1.5b	reciprocal	24	0.9413	0.38	12.02	+0.0065

247 5 Results

248 187 attempt-records were collected across 150 replicate slots (2 models $\times$ up to 3 conditions $\times$ 30
249 replicates); 142 slots kept a usable trace, 8 were residual discards after three attempts. Every cell
250 reached its full 30 replicates.

251 5.1 The dose–response relationship, and the control that overturns it (C1)

252 Across all 85 GARP-violating traces with a computed, independently-verified projection (60 at
253 1.5B, 25 at 3B): Spearman $\rho = 0.729$ ($p < 0.0001$), Pearson $r = 0.821$ ($p < 0.0001$), mean
254 Δ payoff +0.0091 (one-sample $t = 5.41$, $p < 0.00001$), with 70 of 85 traces (82%) improving after
255 repair and 15 (18%) worsening. Figure 2 shows the relationship split by model on shared axes.
256 **Both scales show a significant, positive relationship, and it is stronger at the headroom model:**
257 $\rho = 0.756$ ($p < 0.0001$, $n = 60$) at 1.5B versus $\rho = 0.614$ ($p = 0.0011$, $n = 25$) at 3B. We
258 tested the mechanical confound that larger-dose traces might simply start from worse raw payoff
259 and therefore have more room to improve. The confound is real (dose vs. raw payoff, Pearson
260 $r = -0.41$, $p < 0.0001$) but does not explain the relationship away: the partial correlation of dose
261 against Δ payoff controlling for raw payoff is $r = 0.784$ ($p = 7 \times 10^{-19}$), and an OLS regression of
262 Δ payoff on dose and raw payoff jointly finds dose highly significant ($t = 11.4$, $p < 10^{-16}$) with raw
263 payoff’s own marginal contribution not distinguishable from zero ($p = 0.18$) once dose is included.
264 Full detail in Appendix A.

265 The confound just tested is a confound on the *severity* of the raw violation, not on the specific
266 *geometry* of the payoff function — it says nothing about whether restoring GARP-consistency
267 *specifically*, rather than merely displacing choices toward some interior point, is what buys the payoff
268 gain above. §3.3 builds the control that closes this gap: a GARP-blind null operator, spending the
269 identical L_1 displacement budget as the real repair, that does nothing but shrink every bundle toward
270 the exogenous payoff’s own fixed optimum.

271 **Experiment 1 (original payoff).** The null operator outperforms the real repair by roughly $2.4\times$: mean
272 Δ payoff_{null} = 0.0220 versus mean Δ payoff_{real} = 0.0091, Wilcoxon signed-rank $p = 3.98 \times 10^{-10}$,
273 with the null winning 69 of 85 traces (81%). This holds independently at both models. The partial
274 correlation of dose against Δ payoff_{real}, controlling for the null’s own mechanically-predicted gain,
275 remains positive and significant ($r = 0.574$, $p = 1.13 \times 10^{-8}$) — dose retains real information beyond




Figure 2: Real repair’s Δ payoff (x-axis) against the size-matched null operator’s Δ payoff (y-axis), one point per GARP-violating trace ($n = 85$; circle/blue: qwen2.5:1.5b, triangle/orange: llama3.2:3b). Dashed line is $y = x$; points above it have higher null Δ payoff than real. Symmetric-log axes keep one far-larger-dose trace from compressing the rest. (A) Experiment 1 (original payoff, §3.2). (B) Experiment 2 (corrected payoff, §3.3), each trace’s primary-null Δ payoff averaged over $K = 20$ draws of the random target (± 1 SD error bars); the oracle null is not plotted. Statistics are in §5.1.

276 what the null predicts — but the paired comparison is unambiguous: at the identical displacement
 277 budget, an operator with no knowledge of GARP buys more than twice the payoff the real repair
 278 does.

279 **Experiment 2 (corrected payoff).** The original payoff’s exploit is a single fixed target ($s = 0.5$)
 280 identical across every trace; §3.3 removes this by drawing an independent Cobb–Douglas weight
 281 $\alpha_s \sim \text{Uniform}(0.05, 0.95)$ per trace, repeated across $K = 20$ independent draws for robustness. The
 282 qualitative result is unchanged: the primary, information-fair null still outperforms the real repair
 283 (mean $\Delta\text{payoff}_{\text{null}} = 0.02167$ versus mean $\Delta\text{payoff}_{\text{real}} = 0.00723$, Wilcoxon $p = 1.24 \times 10^{-5}$, win
 284 rate 70.8%), significant in 20 of 20 draws. An oracle null privileged to know the true per-trace target
 285 does better still (mean $\Delta\text{payoff}_{\text{null}} = 0.02224$, $p = 1.35 \times 10^{-9}$, win rate 80.4%, significant in 20/20
 286 draws), reported for completeness rather than as the primary comparison. The partial correlation
 287 of dose with $\Delta\text{payoff}_{\text{real}}$, controlling for the null, attenuates from Experiment 1’s $r = 0.574$ to a
 288 mean $r \approx 0.37$ across the 20 draws, significant in 14–15 of 20 — a real but more modest signal, not
 289 distinguishable from the draw-to-draw noise the 20 draws themselves already show.

290 **No third payoff was attempted, and this is a stated scope decision, not a resource constraint.**
 291 Two independently designed, independently motivated payoffs — one the paper’s original, one built
 292 specifically to remove the first one’s known exploitable structure — converging on the same negative
 293 paired-comparison result is the finding. Redesigning the payoff a third time only because the second
 294 attempt also came back negative would be payoff-shopping: iterating the yardstick until GARP-
 295 restoration happens to look good, exactly the researcher-degree-of-freedom problem an exogenous
 296 payoff was introduced to rule out.

297 **A partial mechanistic note.** Trace extremity (1 — raw payoff under the original fixed payoff, a
 298 stable measure of the raw bundles independent of which payoff later scores them) predicts the null’s
 299 advantage strongly in both experiments (Experiment 1: Pearson $r = 0.679$; Experiment 2 primary
 300 null: mean $r = 0.631$ across 20 draws), more strongly than it predicts the real repair’s own gain
 301 in either. This pattern alone does not cleanly separate “centering” from a more general “moving
 302 an extreme trace toward the interior helps, regardless of exactly where” account, since Experiment
 303 2’s primary null is mechanically the same fixed-center operator as Experiment 1’s. Experiment 2’s
 304 oracle null, which targets a genuinely different, per-trace-varying point, shows essentially the same
 305 extremity-advantage (mean $r = 0.624$) — read as partial, not conclusive, support for an “escape a
 306 bad start” mechanism rather than pure centering, offered as a discussion-level hypothesis rather than
 307 a load-bearing result with its own significance test. Figure 3 illustrates the same idea concretely on

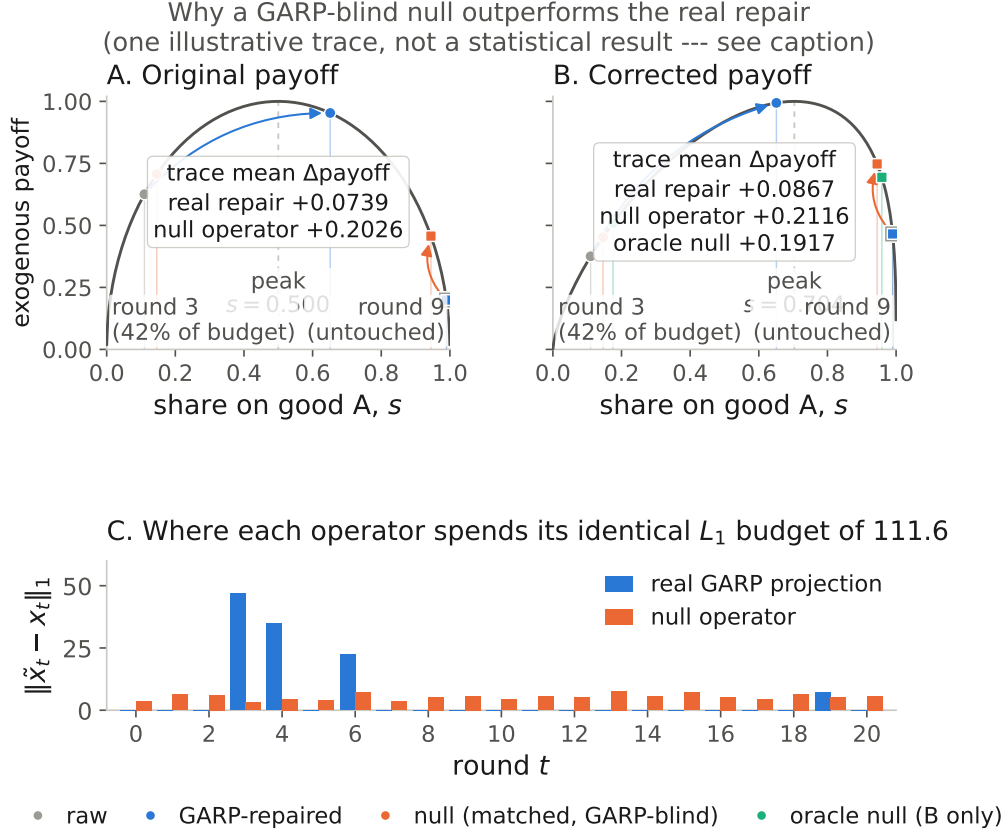


Figure 3: The largest-dose trace (qwen2.5:1.5b, reciprocal, replicate 12; $L_1 = 111.64$, $T = 21$ rounds, hand-audited in Appendix A); every plotted quantity is recomputed by re-solving the projection MILP. With $K = 2$, payoff is a function of expenditure share s . (A) The original payoff, $2\sqrt{s(1-s)}$ (peak at $s = 0.5$). (B) The corrected payoff’s per-trace-target form (peak at this trace’s own draw $\alpha_s = 0.704$). (C) Per-round expenditure share s across all 21 rounds for the raw trace, the real repair, and the null operator. The repair changes 5 of 21 rounds (42% of its L_1 budget on round 3 alone: s 0.11 $\rightarrow$ 0.65, payoff 0.626 $\rightarrow$ 0.953), leaving the other 16 at a near-corner $s = 0.99$; the null spreads the same total displacement over all 21 rounds. Single-trace illustration; the 85-trace comparison is in §5.1.

308 a single trace: the rounds that generate the null’s advantage there are precisely the extreme-share
309 rounds the real repair leaves untouched.

310 **C1 as originally proposed — that restoring GARP-consistency specifically, beyond mere dis-**
311 **placement toward the interior of the budget line, buys more exogenous payoff — is not supported**
312 **under either payoff design.** The raw dose–response relationship reported at the top of this section is
313 real; it is not evidence for the specific mechanism C1 claims.

314 5.2 The reciprocal-framing manipulation: survivorship, and a correction that helps but does 315 not fully resolve it

316 At 3B, reciprocal framing replicates the pilot’s finding on an independently-drawn budget-set design:
317 GARP pass rate collapses from 0.73 (22/30) to 0.39 (11/28), $p = 0.0089$ (BH $p_{\text{BH}} = 0.011$, Holm
318 $p_{\text{Holm}} = 0.062$); CCEI moves much less, reaching only borderline significance (0.9900 vs. 0.9713,
319 t -test $p = 0.0508$) — the same near-1 compression the pilot flagged, and a further case of GARP
320 pass rate being the sensitive instrument for framing/format disruption while CCEI understates it.

At 1.5B, the manipulation this study was designed around does *not* survive proper correction for discard-selection. First-attempt discard under reciprocal framing was 43.3% (13/30), falling to a residual 20.0% (6/30) after the retry protocol — smaller than the pilot’s unretried 52%, but the pattern is unchanged: reciprocal framing produces far more discards than any other condition, including the *other* 1.5B manipulation (multiturn, 0%). Once the surviving 24 sessions are compared, GARP pass rate (0.40 vs. 0.38, $p = 0.85$) and CCEI (0.9522 vs. 0.9413, $p = 0.73$) are indistinguishable between baseline and reciprocal framing. This is no detectable CCEI shift, not a confirmation of the pilot’s own naive estimate: the pilot’s +0.0169 ($p = 0.66$) and the main experiment’s −0.0109 ($p = 0.73$) are two statistically indistinguishable nulls, and reading their sign difference as a reversal overstates what either number supports. Table 2 gives the finer-grained per-attempt breakdown these top-line numbers average over.

5.3 The multiturn/format effect: a large effect in the opposite direction from the literature

Splitting the 25 rounds into separate sequential calls collapses the GARP pass rate from 40% to 10% ($p = 0.0073$) while CCEI barely moves (0.9522 vs. 0.9540, $p = 0.91$; Table 4) — a larger drop than reciprocal framing produced at 1.5B, with *zero discards on either arm*. This survives Benjamini–Hochberg correction ($p_{\text{BH}} = 0.0094$) but not Holm ($p_{\text{Holm}} = 0.058$); we report both. Together with the 3B reciprocal-framing result (§5.2), this is a further case of the same pattern. Against the literature this arm was designed against [Wang et al., 2025a]: in the same domain, on the same model family at a larger scale (Qwen2.5-7B), single-turn elicitation was *more* coherence-degrading than multi-turn; here, at 1.5B, the reverse holds. We do not have an explanation for the reversal; we report the disagreement as a finding in its own right.

Table 4: 1.5B (qwen2.5:1.5b), baseline vs. multiturn elicitation format (25 sequential calls), $n = 30$ per condition, zero discards in either arm. GARP pass rate is reported as count/ n with 95% Wilson score confidence intervals (Pearson χ^2 test, $p = 0.0073$); mean CCEI is reported with 95% confidence intervals computed from the reported mean, SD, and n via the t -distribution ($df = 29$) (two-sample t -test, $p = 0.91$).

Condition	GARP pass	95% CI (Wilson)	mean CCEI [95% CI (t)]
baseline	12/30 (40.0%)	[24.6, 57.7]%	0.9522 [0.9253, 0.9791]
multiturn (25 seq. calls)	3/30 (10.0%)	[3.5, 25.6]%	0.9540 [0.9350, 0.9729]

6 Limitations

CCEI is noisy at 1.5B, and the paper leads with GARP pass rate because of it, not by preference. The 1.5B model’s within-condition CCEI noise (baseline WSCV 14.8%) is comparable to published human test–retest noise on the same index [Nitsch et al., 2022], and $N = 30$ is underpowered for a CCEI-scale effect of the pilot’s own magnitude (minimum detectable effect $\approx 2.5\times$ the pilot’s). A CCEI-power-matched design would need $N \approx 111$ –161; that scale-up was deferred, not executed. This is an open design tradeoff: the pass-rate metric is adequately powered throughout, and every CCEI result carries its own confidence interval.

Single run per condition, one scale point per model. Every (model, condition) cell is one run of the stated protocol, with no repeated full-pipeline runs to report a between-run variance on the headline statistic itself, only the within-run replicate variance in the reported confidence intervals and p -values. The headroom/null-control design isolates the identification strategy from a scale confound but does not trace a dose–response relationship *across* scale: a pilot found headroom and measurement reliability do not coexist anywhere piloted (a third model family was excluded after the multi-model-residency failure of §4).

Two fixed exogenous payoffs, not three. We report two independently designed payoffs (§3.2, §3.3): the first’s single, globally-identical optimum turned out exploitable by a GARP-blind operator; the second removes that property and reproduces the same verdict. A third payoff, run only because the second also came back negative, would be payoff-shopping, not a robustness check (§5.1).

An adverse prior the paper now joins rather than argues against. Three independently published repair attempts moved the wrong way (§2); our reciprocal-framing manipulation at 1.5B (§5.2) is

a fourth, in the narrower sense that it failed to induce the coherence variation it was designed to. The null-operator result of §5.1 is a fifth: under two payoffs, restoring GARP-consistency does not outperform a matched GARP-blind operator. We do not read this as evidence the coherence-competence relationship is negative in general — C2 remains open, not settled — only that this paper’s own attempt to find a positive relationship, net of displacement, did not succeed.

Broader impacts. This paper’s central result argues against imposing GARP-restoration on deployed AI economic agents as a default decision-quality intervention: under two independently designed payoffs, a GARP-blind operator that merely displaces choices toward an interior point outperforms the real repair (§5.1). A practitioner deploying coherence-repair on the strength of a raw dose-response correlation alone would be adopting exactly the confound this paper’s control was built to catch. The reciprocal-framing manipulation this study was designed around similarly turned out, once corrected, to have no measurable effect on the coherence it was thought to disrupt (§5.2). This is a measurement design for asking whether repair helps in a given setting — and, here, a documented instance where it does not — not a recommendation to deploy it universally, nor a claim that no coherence-repair method could ever help. The discard-selection finding (§4.3) generalizes further: any consistency instrument that silently drops disruption-caused failures risks underestimating exactly the manipulations that matter most. We see no elevated misuse risk specific to this work: synthetic tasks, open-weight models, no new dataset, no human subjects.

7 Conclusion

This paper set out to measure whether restoring GARP-consistency in an LLM agent’s own realized choices buys more exogenous payoff than the mechanical fact of displacement alone. Applying Demuynck and Rehbeck [2023]’s minimal-quantity-error GARP repair post hoc to a frozen agent, we find a real, monotone raw dose-response relationship between repair size and payoff gain — but a distance-matched, GARP-blind null operator outperforms the real repair on that same payoff, significantly, at both model scales, under two independently designed payoffs. C1, as originally proposed, is not supported. A partial mechanistic finding — trace extremity predicts the null’s advantage in both experiments — offers a candidate explanation without fully resolving it. What survives is the paper’s identification argument (C2): the coherence-competence relationship remains open, not settled, and this is the first attempt we know of to measure it without confounding coherence with capacity, a preference-judgment outcome, or mere displacement magnitude. Two further, unpredicted findings — a framing manipulation’s apparent effect was a discard-selection artifact rather than a real coherence shift (C3), and an elicitation-format manipulation collapsed GARP pass rate in the opposite direction from a published finding in the same domain and model family at a larger scale — are reported in their own right: a well-controlled negative, an open identification question, and an instrument-validity finding are exactly what this literature should surface by default rather than paper over.

References

- S. N. Afriat. The construction of utility functions from expenditure data. *International Economic Review*, 8(1):67–77, 1967.
- Victor H. Aguiar and Nail Kashaev. GARP-EFM: Improving foundation models with revealed preference structure. *arXiv preprint arXiv:2603.23993*, 2026.
- James Andreoni and John Miller. Giving according to GARP: An experimental test of the consistency of preferences for altruism. *Econometrica*, 70(2):737–753, 2002.
- Isaiah Andrews. Revealed rationality: Label-free evaluation and regularization from representation theorems. *arXiv preprint arXiv:2608.05015*, 2026.
- Stephen G. Bronars. The power of nonparametric tests of preference maximization. *Econometrica*, 55(3):693–698, 1987.
- Joy Buchanan and Joshua Foster. The innate economic preferences of language models. *arXiv preprint arXiv:2607.26288*, 2026.

Alina Chadwick, Anson Kahng, and Jens Kipper. Dutch books and money pumps: rectifying vulnerabilities in LLMs through rationality. In *Proceedings of the 4th International Conference on Human and Artificial Rationality (HAR)*, 2025.

Yiting Chen, Tracy Xiao Liu, You Shan, and Songfa Zhong. The emergence of economic rationality of GPT. *Proceedings of the National Academy of Sciences*, 120(51):e2316205120, 2023.

Syngjoo Choi, Raymond Fisman, Douglas Gale, and Shachar Kariv. Consistency and heterogeneity of individual behavior under uncertainty. *American Economic Review*, 97(5):1921–1938, 2007.

Syngjoo Choi, Shachar Kariv, Wieland Müller, and Dan Silverman. Who is (more) rational? *American Economic Review*, 104(6):1518–1550, 2014.

Thomas R. Cook, Sophia Kazinnik, Zach Modig, and Nathan M. Palmer. What do LLMs want? Technical Report FEDS 2026-006, Federal Reserve Board, 2026.

Mark Dean and Daniel Martin. Measuring rationality with the minimum cost of revealed preference violations. *Review of Economics and Statistics*, 98(3):524–534, 2016.

Thomas Demuynck and John Rehbeck. Computing revealed preference goodness-of-fit measures with integer programming. *Economic Theory*, 76(4):1175–1195, 2023.

M. Houtman and J. A. Maks. Determining all maximal data subsets consistent with revealed preference. *Kwantitatieve Methoden*, 19:89–104, 1985.

Huang, Li, Zhao, and Li. Transitivity meets cyclicity: Explicit preference decomposition for dynamic llm alignment. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2605.17342.

Q. Huangfu and J. A. J. Hall. Parallelizing the dual revised simplex method. *Mathematical Programming Computation*, 10(1):119–142, 2018.

Felix J. Nitsch, Luca M. Lüpken, Nils Lüscho, and Tobias Kalenscher. On the reliability of individual economic rationality measurements. *Proceedings of the National Academy of Sciences*, 119(31):e2202070119, 2022.

Shumiao Ouyang, Hayong Yun, and Xingjian Zheng. AI as decision-maker: Ethics and risk preferences of LLMs. *arXiv preprint arXiv:2406.01168*, 2024.

Mingxuan Sun. CONSISTRE: A unified consistency-aware framework for document-level relation extraction with large language models. *arXiv preprint arXiv:2607.24312*, 2026.

Hal R. Varian. The nonparametric approach to demand analysis. *Econometrica*, 50(4):945–973, 1982.

Wang, Yao, Zhang, Gai, Liu, and Zhong. When experimental economics meets large language models: Tactics with evidence. *arXiv preprint arXiv:2505.21371*, 2025a.

Wang, Zhan, et al. Trustroboreward: Preference-ordered isotonic score editing for multi-paradigm robot reward modeling. *arXiv preprint arXiv:2608.08491*, 2026.

Yidong Wang, Yunze Song, Tingyuan Zhu, Xuanwang Zhang, Zhuohao Yu, Hao Chen, Chiyu Song, Qiufeng Wang, Cunxiang Wang, Zhen Wu, Xinyu Dai, Yue Zhang, Wei Ye, and Shikun Zhang. Trustjudge: Inconsistencies of LLM-as-a-judge and how to alleviate them. *arXiv preprint arXiv:2509.21117*, 2025b.

Stephen Warshall. A theorem on Boolean matrices. *Journal of the ACM*, 9(1):11–12, 1962.

Kaiser Yamin, J. Tang, S. Cortes-Gomez, A. Sharma, Eric Horvitz, and Bryan Wilder. When agents say one thing and do another: Validating elicited beliefs from LLMs. *arXiv preprint arXiv:2602.06286*, 2026.

Yifan Zeng, Ojas Tendolkar, Raymond Baartmans, Qingyun Wu, Lizhong Chen, and Huazheng Wang. LLM-RankFusion: Mitigating intrinsic inconsistency in LLM-based ranking. *arXiv preprint arXiv:2406.00231*, 2024.

- 457 Zhang, Zhang, Wu, Xu, and Gu. Beyond bradley-terry models: A general preference model for
 458 language model alignment. In *International Conference on Machine Learning (ICML)*, 2025.
 459 arXiv:2410.02197.
- 460 Jian-Qiao Zhu, Haijiang Yan, and Thomas L. Griffiths. Recovering event probabilities from large
 461 language model embeddings via axiomatic constraints. *arXiv preprint arXiv:2505.07883*, 2025.

462 A Payoff exogeneity audit

463 Full detail of the four-part adversarial audit summarized in §3.2.

464 **(1) No shared derived quantity beyond the exogenous (p, I) .** The payoff implementation has no
 465 dependency on the projection implementation; the optimal bundle x_t^* is a function of (p_t, I_t) alone
 466 and never references x_t or $\tilde{x}_t$. The one near-miss is that the projection’s feasibility incumbent is a
 467 Cobb–Douglas demand *share-fitted to the agent’s own observed data* — a plausible leakage channel
 468 on first read — but this quantity is never passed to the underlying MILP solver as an actual incumbent
 469 (the solver used has no such interface); it is used only to log an upper-bound sanity ceiling on the
 470 reported distance, never to steer the solution.

471 **(2) The payoff’s optimum is independent of revealed preference, checked empirically as well
 472 as by construction.** Comparing raw (pre-repair) payoff between the 57 traces that were already
 473 GARP-consistent and the 85 that were not: mean 0.7872 vs. 0.7899, Welch $t = -0.07$, $p = 0.94$.
 474 GARP-consistency status alone does not predict payoff, ruling out a population-level mechanical link
 475 between coherence and payoff score.

476 **(3) Hand-derived mechanism check.** Five traces spanning the full dose range (0.17 to 111.64)
 477 were inspected directly, comparing the direction of the actual projection movement $(\tilde{x} - x)$ against
 478 the direction toward the exogenous optimum $(x^* - x)$ via cosine similarity. Similarities were low
 479 (0.16–0.31) on the four traces where payoff improved after repair, and *negative* (−0.35) on the single
 480 trace where payoff worsened — the sign of the outcome tracks a geometric quantity the projection
 481 objective never references, evidence against the mechanism being a disguised optimization toward
 482 the payoff target.

483 **(4) The severity-confound check.** Larger-dose traces do start from worse raw payoff (Pearson
 484 $r = -0.41$, $p < 0.0001$), and worse-starting traces do have more room to improve ($r = -0.41$
 485 between raw payoff and Δ payoff) — a genuine ceiling-effect confound. It does not explain the
 486 dose–response relationship away: the partial correlation of dose against Δ payoff controlling for raw
 487 payoff is $r = 0.784$ ($p = 7 \times 10^{-19}$, attenuated only slightly from the unconditional $r = 0.821$), and
 488 in a joint OLS regression the dose coefficient is highly significant ($t = 11.43$, $p < 10^{-16}$) while raw
 489 payoff’s own coefficient is not ($t = -1.36$, $p = 0.18$).

490 B Full MILP formulation and method-implementation detail

491 Full detail of the projection’s exact formulation and implementation commitments, summarized in
 492 §3.1.

493 **Formulation.** The naive formulation parameterizes the projection with Afriat multipliers, which
 494 makes the resulting constraints bilinear in the unknowns once $\tilde{x}$ is itself a decision variable — fixing
 495 a candidate preference ordering does not remove the bilinearity, since a price-weighted multiplier
 496 term remains a product of two unknowns regardless of ordering. We instead use the multiplier-free
 497 ordinal characterization of Demuynck and Rehbeck [2023], under which every constraint is linear
 498 jointly in the perturbed bundles, a set of ordinal utility levels $u_t \in [0, 1]$ that carry no multipliers, and
 499 binary comparison indicators $U_{t,v} \in \{0, 1\}$ for $t \neq v$, with no outer search over preference orderings
 500 and no alternating scheme.

501 **Three method commitments, each carrying an explicit risk. Budget exhaustion is imposed as an
 502 equality**, $p_t \cdot \tilde{x}_t = I_t$ for every observation (income fixed at 100 throughout); this makes the big- M
 503 constant in the ordering constraints computable a priori as $\alpha > \max_t I_t$. **A fixed strict-preference
 504 margin $\gamma > 0$** , set to $10^{-4} \cdot \min_t I_t$, converts an unattained infimum into an attained minimum: the
 505 GARP-consistent set is not closed, so an unregularized projection can read a distance of exactly zero
 506 on genuinely violating data. We verified the reported distance is stable across $\gamma \in \{10^{-2}, \dots, 10^{-6}\}$.

507 **Every returned $\tilde{x}$ is verified independently** via the same combinatorial Warshall-closure [Warshall,
 508 1962] GARP check used everywhere else, never trusted from solver optimality status alone — all
 509 85 GARP-violating traces’ projections were independently re-verified GARP-consistent, MIP gap
 510 $\leq 8.1 \times 10^{-5}$ on every solve.

511 **Objective.** $\min \sum_t \sum_k w_{t,k} d_{t,k}$ subject to $\tilde{x}_{t,k} - x_{t,k} \leq d_{t,k}$ and $x_{t,k} - \tilde{x}_{t,k} \leq d_{t,k}$, with uniform
 512 weights $w = 1$.

513 **What can be guaranteed, and what cannot.** Against the closest published structural analogue, a
 514 proved-optimal projection onto a convex monotone cone [Wang et al., 2026]: because the GARP-
 515 consistent set is a union of polyhedra rather than a single convex set, the non-expansiveness that
 516 licenses a weakly-closer-to-ground-truth guarantee there has no analogue here — the prior operator’s
 517 cone is fixed by an ordering supplied as input, whereas a GARP repair must decide which revealed-
 518 preference comparisons to give up, absorbed into the binary comparison indicators above rather
 519 than eliminated, which is why no distance-minimization guarantee analogous to the convex case is
 520 claimed.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and §1 state exactly what is claimed (a measured dose–response relationship, an instrument-validity finding, and a replication of an independently published format effect) and what is not (novelty of repair, of projection, or of the sign of the effect); §6 states the design’s scope limits explicitly.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: §6 (Limitations) discusses CCEI’s noise and power at 1.5B, the single-run-per-condition design, the two-model-family scope, the single fixed payoff weighting, and the adverse prior from three prior negative results.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper’s contribution is a computational method (a MILP formulation for GARP-repair) building on a cited external representation theorem [Demuyne and Rehbeck, 2023]; it does not state or prove new theorems of its own.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §4-5 specify the exact models, budget-set generator and seeding scheme, conditions, power calculation, retry protocol, and the full MILP formulation (§3) needed to reproduce the main results; the exogenous payoff is a closed-form function stated in full.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: An anonymized code and data package is included as supplementary material (prompts, seeds, the parser/validity rule, and raw per-trace logs for the main experiment, the null-operator control, the corrected-payoff robustness check, and the pilot), with a README giving the exact commands to regenerate every reported statistic from the included raw logs without re-running any model inference. Both LLMs used are publicly available open-weight models, so the full pipeline is also reproducible from scratch.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: §4 specifies model identities, all three conditions' exact framing, the budget-set generator and its rejection rule, seed scheme, replicate count and power calculation, and the discard/retry protocol.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Every reported comparison in §5 carries a p -value, and Table 1 and the dose-response statistics report means alongside sample sizes; CCEI comparisons additionally report standard deviations and, in §6, confidence-interval-based power limitations.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All inference was run locally on a single consumer CPU-class machine (no GPU, no commercial API calls), with the full $N = 30$ design measured at approximately 2.1 hours wall-clock (§4.1); every MILP projection solves in under 5 seconds with MIP gap $\leq 8.1 \times 10^{-5}$.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The study elicits only synthetic, non-sensitive budget-allocation choices from open-weight language models; it involves no human subjects, no personal data, and no scraped content.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: See the Broader impacts paragraph preceding §7: the paper’s own finding that 18% of repairs worsen the exogenous payoff is itself a caution against treating coherence-repair as a costless, universally beneficial deployment default.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no model or dataset; both LLMs used are existing, publicly released open-weight models used unmodified as measurement subjects.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The two LLMs used are named with version in §4.1. Qwen2.5-1.5B-Instruct (Alibaba Cloud) is released under the Apache License 2.0. Llama-3.2-3B-Instruct (Meta Platforms, Inc.) is released under the Llama 3.2 Community License Agreement, a custom commercial license. Both are publicly available open-weight models used here for research purposes consistent with their license terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: This submission does not release a new dataset or model asset; the generated budget-set choice logs are experimental records, not a packaged new asset.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The study involves no crowdsourcing and no human subjects; all choices are elicited from language models.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The study involves no human subjects and required no IRB review.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: The two LLM agents under study are the paper's core experimental subject, not a writing aid: §4.1 names both models and their role as the choice-making agents whose revealed preferences are elicited, projected, and scored.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.